

# CS-GY 6923 MACHINE LEARNING

*Professor: Dr. Raman Kannan*

## **HW1: First Modeling Assignment**

Remon Roshdy

# Contents

1	Overview of the Data set	2
2	Loading the Data	2
3	Exploratory Data Analysis	3
4	Random Forest Model and Variable Importance	10
5	Logistic Regression Classification Using Top 10 Features	12
6	Null and Residual Deviance	15
7	Eliminating Correlated Features	15
8	Logistic Regression Model Using Uncorrelated Features	20
9	Logistic Regression Model Performance Metrics	23
10	Naive Bayes Classifier	29
11	kNN Classifier	35
12	Dealing with the Class Imbalance Issue	38
13	Dimensionality Reduction Using PCA	44

# 1 Overview of the Data set

The data set is called "Smartphone-Based Recognition of Human Activities and Postural Transitions" and was taken from the UCI repository:

<http://archive.ics.uci.edu/ml/datasets/Smartphone-Based+Recognition+of+Human+Activities+and+Postural+Transitions>

It is an activity recognition data set built from the recordings of 30 subjects performing basic activities and postural transitions while carrying a waist-mounted smartphone with embedded inertial sensors. The features are composed of parameters calculated from the signals obtained through the inertial sensors, and the activity is recorded as the output variable to be predicted.

The data set consists of 561 features (independent variables), 1 dependent variable and 7767 observations. The dependent variable has categorical integer values ranging from 1 to 12; hence, this is a classification problem. And since the number of classes is larger than two, it is a multi-class classification. Before loading the data file into R, labels were assigned to the variables, and the file was saved as a csv file. The label "activity" was allocated to the class variable, and the independent variables were assigned with numerical labels ranging from 1 to 561 for better representation and easier manipulation of the large number of features.

## 2 Loading the Data

The data set, containing all the variables and labels, was saved as a csv file named 'train\_num.csv' and it was loaded into a variable called data.

---

```
1 > loadData=function(csvfile) { read.csv(csvfile,head=T,sep=',',stringsAsFactors=F
    ) } # Function to load the data
2 > data=loadData('train_num.csv') # Load the data
```

---

We then perform few steps to inspect the data to make sure that it was loaded properly. The dimensions of the data frame was found to be  $7767 \times 562$ , which corresponds to 7767 rows (number of observations) and 562 columns (number of variables). The number of columns includes the 561 features plus 1 column for the class variable. We also check that the labels are assigned correctly using the `names(data)` command. Since we labeled the independent variables with numerical values, R automatically inserts an "X" prior to the label as shown in the output. We can also see that the last column is the class variable, and its name is "activity". All the features (independent variables) columns have continuous numerical values, so they are non-categorical.

---

```

1 > head(data)[,1:6] # Display the first few entries of the data
2           X1           X2           X3           X4           X5           X6
3 1 0.04357967 -0.005970221 -0.03505434 -0.9953812 -0.9883659 -0.9373820
4 2 0.03948004 -0.002131276 -0.02906736 -0.9983480 -0.9829449 -0.9712729
5 3 0.03997778 -0.005152716 -0.02265071 -0.9954821 -0.9773138 -0.9847595
6 4 0.03978456 -0.011808778 -0.02891578 -0.9961941 -0.9885686 -0.9932556
7 5 0.03875814 -0.002288533 -0.02386289 -0.9982413 -0.9867741 -0.9931155
8 6 0.03898801 0.004108852 -0.01734027 -0.9974376 -0.9934854 -0.9966920
9 > dim(data) # Check the number of rows and columns of the data
10 [1] 7767 562
11 > names(data) # Check the names of each column (features and class variable)
12 [1] "X1"      "X2"      "X3"      "X4"      "X5"      "X6"
13 [7] "X7"      "X8"      "X9"      "X10"     "X11"     "X12"
14 [13] "X13"     "X14"     "X15"     "X16"     "X17"     "X18"
15 [19] "X19"     "X20"     "X21"     "X22"     "X23"     "X24"
16 [25] "X25"     "X26"     "X27"     "X28"     "X29"     "X30"
17 [31] "X31"     "X32"     "X33"     "X34"     "X35"     "X36"
18 [37] "X37"     "X38"     "X39"     "X40"     "X41"     "X42"
19 [43] "X43"     "X44"     "X45"     "X46"     "X47"     "X48"
20 [49] "X49"     "X50"     "X51"     "X52"     "X53"     "X54"
21 [55] "X55"     "X56"     "X57"     "X58"     "X59"     "X60"
22 [61] "X61"     "X62"     "X63"     "X64"     "X65"     "X66"
23 [67] "X67"     "X68"     "X69"     "X70"     "X71"     "X72"
24 ...
25 [547] "X547"     "X548"     "X549"     "X550"     "X551"     "X552"
26 [553] "X553"     "X554"     "X555"     "X556"     "X557"     "X558"
27 [559] "X559"     "X560"     "X561"     "activity"
28 > length(names(data)) # Double check that the length of "names" matches the
    number of columns
29 [1] 562
30 > which(names(data)=='activity') # Position of the class variable column
31 [1] 562

```

---

### 3 Exploratory Data Analysis

Before we can train a model, we need to understand the data and perform some analyses in order to ensure that it is consistent with the classifier to be used. In Section 2, we saw that we have 7767 observations and 561 features; hence, this is considered to be a large data set as it is, which requires heavy computations and not all classifiers might be able to handle depending on the resources available. Let's first find the spectral count of the class variable "activity", which will give us the number of classes, class names, and the number of observations for each class.

---

```

1 > # Find spectral count
2 > TBL=table(data$activity)
3 > TBL
4      1      2      3      4      5      6      7      8      9     10     11     12
5 1226 1073  987 1293 1423 1413   47   23   75   60   90   57
6 > names(TBL)
7 [1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10" "11" "12"
8 > as.numeric(TBL)
9 [1] 1226 1073  987 1293 1423 1413   47   23   75   60   90   57
10 > # Determine if binary or multi-class classification
11 > print(ifelse(length(TBL)==2, "Binary Classification", "MultiClass
      Classification"))
12
13 [1] "MultiClass Classification"

```

---

The output of the `table(data$activity)` command shows that we have 12 classes, each named with an integer number from "1" to "12". And as mentioned earlier, since the number of classes are larger than 2, we are dealing with Multi-class classification. Moreover, just by looking at the number of observations for each class label, we can tell that there might be a class imbalance issue. Therefore, we run a check to determine if the number of observations of one class to another is less than 1:6, and in that case, the class would be imbalanced. The results show that we do have class imbalance, and the imbalanced classes are "7", "8", "9", "10", "11", and "12". Further analysis of the imbalance issue, and the process to deal with it, will be discussed at a later stage.

---

```

1 > # Determine if imbalanced. Criteria: class count 1:6
2 > if(any(TBL<(nrow(data)*1/12*1/6) | any(TBL>(nrow(data)*1/12*6)))) {
3 +   imbalanced_classes=names(TBL)[TBL<(nrow(data)*1/12*1/6) | TBL>(nrow(data)*1/12
4 +     *6)]
5 +   print("Imbalanced Classes:")
6 +   print(imbalanced_classes)
7 + } else print("Dataset is Balanced")
8
9 [1] "Imbalanced Classes:"
10 [1] "7" "8" "9" "10" "11" "12"

```

---

Next, we need to check if standardization or normalization of the data is required. The standard deviation and mean of each feature column were calculated, and the results are plotted in Fig. 1. The figure shows that the means are not equal to 0 and the standard deviations are not equal to 1, which implies that the data is not standardized. Nevertheless, the description of the data set mentioned that the data was normalized and bounded between [-1 1]; therefore, to verify this information, we calculate the ranges of each feature column and plot them in Fig. 2, which indicates that the data is normalized within [-1 1] and, thus, no scaling is needed.


---

```

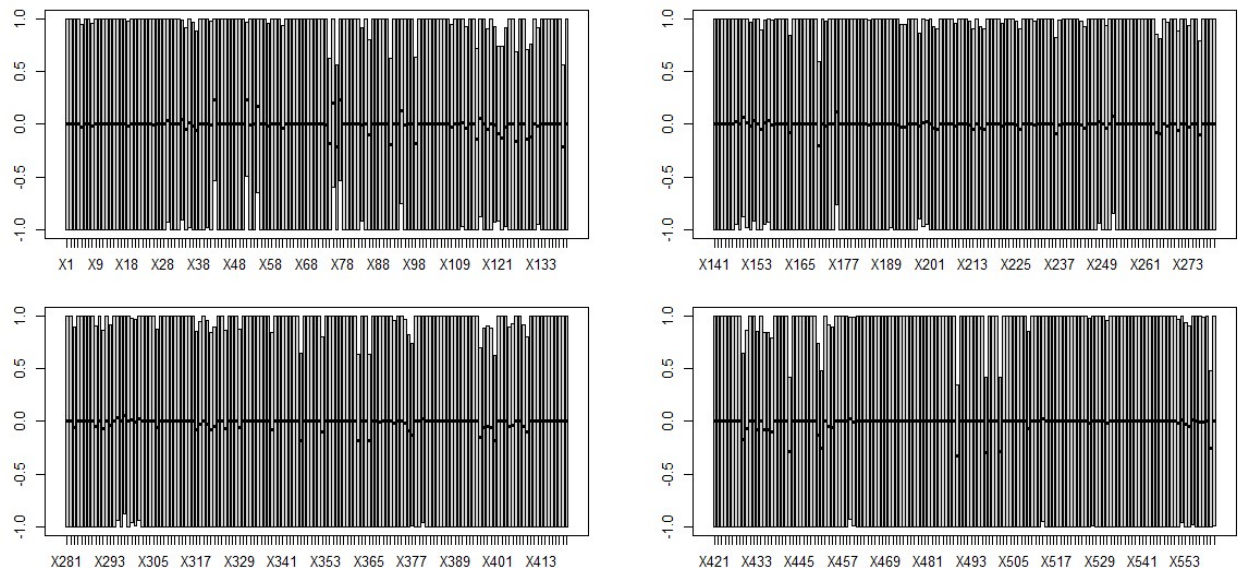
1 > # Check if scaling and standardization are needed
2 > std_dev_data=round(apply(data[, -562], 2, sd), 4) # Standard deviation by column
3 > mean_data=round(apply(data[, -562], 2, mean), 4) # Mean for each column
4 > par(mfrow=c(2, 1), mar=c(2, 4, 2, 2)) # plot the standard deviation and mean of each
    feature column
5 > plot(std_dev_data)
6 > plot(mean_data)
7 > par(mfrow=c(1, 1))
8 > ranges=apply(data[, -which(names(data)=='activity')], 2, range) # min/max for each
    column
9 > par(mfrow=c(2, 2)) # plot the ranges
10 > boxplot(ranges[, 1:140])
11 > boxplot(ranges[, 141:280])
12 > boxplot(ranges[, 281:420])
13 > boxplot(ranges[, 421:561])
14 > par(mfrow=c(1, 1))

```

---



**Figure 1:** Calculated standard deviations (top figure) and means (bottom figures) of each feature column.



**Figure 2:** Boxplots showing the ranges of each feature column. Due to the large number of features, they were split into 4 figures, with each figure containing a subset of the total number of feature columns.

Thereafter, we inspect the data for constant and correlated predictors (features). Constant predictors are not required since they do not add any value in training the model, so they shall be eliminated. Similarly, if two predictors are highly correlated, then collinearity exists and, in this case, we can keep one of them and discard the other since only one of those predictors provides sufficient information to have a proper model fit. At this stage, we will only check if correlated predictors exist, whereas eliminating them will be performed at a later stage after training the first model in order to study the effect of discarding those correlated features. The output of the `pairs(data[,1:8])` command is provided in Fig. 3, which shows the scatter plots between each pair of the first 8 features. By visual inspection, we can tell that there is a strong correlation between features X4 and X7, and features X5 and X8. Using the command `cor(data[,1:561])`, we obtain the correlation matrix which includes the correlation coefficients between each pair of predictors. A correlation coefficient of 0 indicates that the predictors are uncorrelated, and a value of 1 indicates that they are perfectly correlated. Accordingly, the criteria to identify two predictors to be correlated is to have a correlation coefficient larger than 0.5 between those predictors. The results of running the R code below indicate that we have 516 correlated predictors out of the total 561 predictors in our data set.

---

```

1 > # Eliminate constant predictors
2 > const_pred=unlist(lapply(1:561,FUN=function(x) {
3 +   TBL=table(data[[x]])
4 +   ifelse(length(names(TBL))<2,-1*x,x)})) # Finds for each feature column if the
      number of distinct values are less than 2 (values of feature is constant)
5 > print(ifelse(any(const_pred<0),"Constant Predictors Exist","No Constant
      Predictors"))
6 > data=data[,const_pred>0]
7
8 [1] "No Constant Predictors"

```

---


```

1 > # Determine if collinearity exists (correlation between predictors). Criteria:
      correlation coefficient > 0.5
2 > pairs(data[,1:8]) # Plot the predictors one vs the other (first 8 predictors
      only)
3 > cordata=cor(data[,1:561]) # Correlation matrix (correlation coefficients)
4 > print(ifelse(any(abs(cordata[cordata!=1])>0.5),"Correlated Predictors Exist","
      No Correlated Predictors"))
5 > cor_index=which(abs(cordata)>0.5 & abs(cordata)!=1, arr.ind = T) # Get the
      indices where the correlation coefficient > 0.5 in the correlation matrix
6 > cor_index=cor_index[!duplicated(cbind(pmax(cor_index[,1], cor_index[,2]), pmin(
      cor_index[,1], cor_index[,2]))),] # Remove duplicates, i.e. retain one of
      the two correlated predictors
7 > tbl_cor_index=table(cor_index[,1])
8 > cor_index_num=length(tbl_cor_index) # Number of correlated predictors
9 > print(paste("Number of correlated predictors =",cor_index_num))
10
11 [1] "Correlated Predictors Exist"
12 [1] "Number of correlated predictors = 516"

```

---





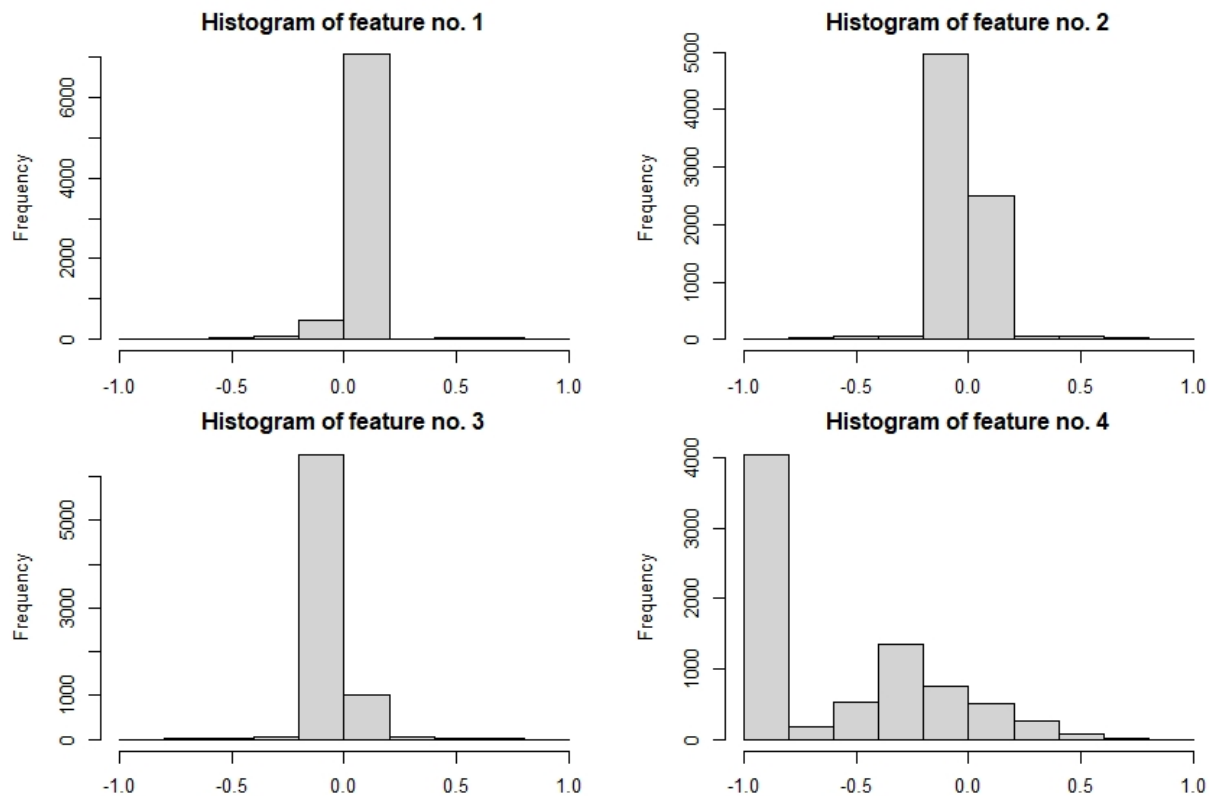
**Figure 3:** Scatter plots generated by the `pairs(data[,1:8])` command.

The next step is to identify the outliers in the data set. Detecting outliers was based on the 3-sigma approach, meaning that any data point or observation that falls outside three standard deviations from the mean, in each feature column, is considered to be an outlier. Histograms of the first four features are plotted and shown in Fig. 4. Here, the values that are away from the main (high frequency) observations are potential outliers. Running the R code over all feature columns, we identified 4714 out of 7767 observations as outliers. However, this identification process was based on each feature column separately; therefore, declaring an observation to be an outlier by this approach is not proper. A better way is to use a generated model and take into consideration the feature's significance and collectively consider the important features in determining if an observation is an outlier or not. Consequently, dealing with outliers will be delayed to a later stage.

```

1 > # Finding outliers
2 > par(mfrow=c(2,2)) # plot histograms of the first 4 features
3 > for(i in 1:4) {hist(data[,i],main=paste("Histogram of feature no.", i))}
4 > par(mfrow=c(1,1))
5 > outliers_row=c() # Loop over the feature columns
6 > for(i in 1:561) {
7 +   data_mean=mean(data[,i]) # Mean of the data in feature column i
8 +   data_sd=sd(data[,i]) # Standard deviation of the data in feature column i
9 +   low_cutoff=data_mean-3*data_sd # Lower cutoff value
10 +   upper_cutoff=data_mean+3*data_sd # Upper cutoff value
11 +   outliers_idx=which(data[,i]<low_cutoff | data[,i]>upper_cutoff) # Get the
      location of outliers in feature column i
12 +   outliers_row=c(outliers_row,outliers_idx) # Append the outlier indeces to
      those of the previous feature columns
13 + }
14 > outliers_row=unique(outliers_row) # Remove duplicated row indeces
15 > print(paste("Number of Outliers =",length(outliers_row)))
16
17 [1] "Number of Outliers = 4714"

```



**Figure 4:** Histograms of the data in the first 4 feature columns.

Now, we need to prepare the data before creating the first classification model. This is done by splitting the data set into a training set and a testing set. To this end, the data was shuffled and 70% of the total number of observations were randomly sampled and stored as the training set, while the remaining 30% were stored as the testing set. After splitting, we need to ensure that the class distributions between the training set, testing set, and original full data set are all somewhat similar, so they're not overly impacted by the sampling process. This is verified by the output of the last three lines in the R code below. Also, it is worth noting that, since this process involved randomization, we should set a seed value (`set.seed(43)`) so that the results are reproducible.

---

```

1 > # Split the data set into train and test data sets
2 > set.seed(43)
3 > randomized=data[sample(1:nrow(data),nrow(data)),] # Shuffle
4 > tridx=sample(1:nrow(data),0.7*nrow(data),replace=F) # Get indices for 70% of
      the total number of samples
5 > trdf=randomized[tridx,] # Define training data set
6 > tstdf=randomized[-tridx,] # Define testing data set
7 > table(data$activity)/nrow(data) # Check if class distribution is similar
8           1           2           3           4           5           6
9 0.157847303 0.138148577 0.127076091 0.166473542 0.183211021 0.181923523
10          7           8           9          10          11          12
11 0.006051242 0.002961246 0.009656238 0.007724990 0.011587486 0.007338741
12
13 > table(trdf$activity)/nrow(trdf)
14           1           2           3           4           5           6
15 0.153973510 0.138337013 0.125091979 0.164091244 0.185982340 0.184878587
16          7           8           9          10          11          12
17 0.005886681 0.002575423 0.009565857 0.007542311 0.013061074 0.009013981
18
19 > table(tstdf$activity)/nrow(tstdf)
20           1           2           3           4           5           6
21 0.166881167 0.137709138 0.131703132 0.172029172 0.176748177 0.175032175
22          7           8           9          10          11          12
23 0.006435006 0.003861004 0.009867010 0.008151008 0.008151008 0.003432003

```

---

## 4 Random Forest Model and Variable Importance

The first objective is to obtain the ten most important features. This can be performed by training a classification model based on the random forest algorithm, using all the features in the data set, and then use that model to compute the variable importance for each feature. The `varImp()` command from `caret` package is used to obtain the variable importance from the random forest model. Fig. 5 shows the computed variable importance for the top 35 features. The features were sorted in descending order according to their variable

importance value, and the top 10 features were extracted and are shown in the output of the R code below.

---

```

1 > # Create a model with random forest and compute variable importance
2 > require(randomForest)
3 > require(caret)
4 > formstr="activity~." # Formula argument with all features
5 > trdf_RF=trdf
6 > trdf_RF$activity=as.factor(trdf_RF$activity) # Makes it run as classification
7 > RF_model=randomForest(activity~.,trdf_RF) # Train the model
8 > RF_model_VI=varImp(RF_model) # Compute Variable Importance
9 > varImpPlot(RF_model,main='Variable Importance') # Plot Variable Importance
10 > RF_model_VI=RF_model_VI[order(RF_model_VI,decreasing=TRUE),,drop=FALSE] # Sort
    from highest to lowest
11 > top10=row.names(RF_model_VI[1:10,,drop=FALSE]) # Top ten features
12 > print(paste("Top ten features are =",paste(top10,collapse=", ")))
13
14 [1] "Top ten features are = X50, X42, X57, X41, X560, X53, X559, X51, X54, X58"

```

---




Figure 5: Variable importance for top 35 features.

## 5 Logistic Regression Classification Using Top 10 Features

Now, we need to generate a logistic regression classifier using the top 10 features as obtained in Section 4, and then compute its classification accuracy as a performance metric. Although the `glm()` function in R can train a logistic regression model, it is not suitable in our case since it can only handle binomial classification. Therefore, we will be using the `multinom()` function from `nnet` package, which is able to train a multinomial logistic regression classifier. The formula argument is defined such that the model will only consider the top 10 features for training. By running the `multinom()` command, we obtained the classifier and used it to perform predictions using the training set (learning phase) and testing set (generalization phase). The confusion matrix was generated and the classification accuracy was extracted in each phase, as shown in the output below. The accuracy in the learning phase was found to be 73.3%. Although the learning phase accuracy is not high enough ( $< 80\%$ ), it does indicate that the model is capable of learning since it is better than random guessing (8.3%). The generalization phase accuracy was found to be 70.5%, which is lower by 3.8% from the learning phase accuracy. Since the drop in accuracy is not larger 25%, we can say that the model is not over-fitting.

---

```
1 > # Create a logistic regression model using the top 10 features
2 > require(nnet)
3 > formstr_top10=paste("activity~",paste(top10,collapse="+")) # Formula argument
   with top 10 features only
4 > mn_model1=multinom(formstr_top10,trdf,maxit=1000) # Train the model
```

---

```
1 > # Predict using the train data (Learning Phase)
2 > mn_pred1_tr=predict(mn_model1,trdf[,which(names(trdf)=="activity")], type="
   class")
3 > mn_cfm1_tr=confusionMatrix(table(trdf[,which(names(trdf)=="activity")], mn_pred
   1_tr)) # Confusion Matrix for train data
4 > print("Learning Phase Confusion Matrix")
5 > mn_cfm1_tr
6 > mn_acc1_tr=round(mn_cfm1_tr$overall[['Accuracy']],4) # Accuracy of predictions
   with train data
7 > print(paste("Accuracy of train data predictions with top 10 features =",mn_acc1
   _tr))
8
```

---

```

9 [1] "Learning Phase Confusion Matrix"
10 Confusion Matrix and Statistics
11
12     mn_pred1_tr
13      1  2  3  4  5  6  7  8  9 10 11 12
14 1  399 80 61 31 266  0  0  0  0  0  0
15 2   73 487 132 23 35  0  0  0  0  2  0
16 3  134 66 435 33 12  0  0  0  0  0  0
17 4    5  3 12 722 136  0  0  4  8  0  2
18 5   42 41 18 86 820  0  1  1  1  0  1
19 6    0  0  0  0  0 994  0  0  0  6  0  5
20 7    0 12  8  1  0  0  8  0  0  0  3  0
21 8    0  0  0  2  0  0  2  9  0  0  1  0
22 9    0  0  1  7  0  0  1  0 29  0  7  7
23 10   0  0  0  0  0  4  0  0  4 19  2 12
24 11   0  7  1  4  0  1  2  3  6  2 43  2
25 12   0  0  0  0  0 11  0  0  4 11  2 21
26
27 Overall Statistics
28
29             Accuracy : 0.7333
30             95% CI : (0.7213, 0.745)
31     No Information Rate : 0.2334
32     P-Value [Acc > NIR] : < 2.2e-16
33
34             Kappa : 0.6835
35
36 [1] "Accuracy of train data predictions with top 10 features = 0.7333"

```

---

```

1 > # Predict using the test data (Generalization Phase)
2 > mn_pred1_tst=predict(mn_model1,tstdf[,-which(names(tstdf=="activity"))], type="
      class")
3 > mn_cfm1_tst=confusionMatrix(table(tstdf[,which(names(tstdf=="activity"))],
      mn_pred1_tst)) # Confusion Matrix for test data
4 > print("Generalization Phase Confusion Matrix")
5 > mn_cfm1_tst
6 > mn_acc1_tst=round(mn_cfm1_tst$overall[['Accuracy']],4) # Accuracy of
      predictions with test data
7 > print(paste("Accuracy of test data predictions with top 10 features =",mn_acc1
      _tst))
8

```

```

9 [1] "Generalization Phase Confusion Matrix"
10 Confusion Matrix and Statistics
11
12     mn_pred1_tst
13      1  2  3  4  5  6  7  8  9 10 11 12
14 1  187 42 26 16 118 0 0 0 0 0 0
15 2   30 200 72 8 10 0 0 0 0 0 1
16 3   64 41 177 19 6 0 0 0 0 0 0
17 4    3 1 7 300 84 0 0 0 4 0 2
18 5   20 18 6 36 330 0 1 0 0 0 1
19 6    0 0 0 0 0 0 405 0 0 0 1 1
20 7    0 2 6 1 0 0 6 0 0 0 0 0
21 8    0 0 0 2 0 0 1 4 0 0 2 0
22 9    0 0 2 1 0 0 0 0 14 1 4 1
23 10   0 0 0 0 0 6 0 0 0 5 2 6
24 11   0 2 1 0 0 0 0 1 3 0 11 1
25 12   0 0 0 0 0 1 0 0 1 1 0 5
26
27 Overall Statistics
28
29             Accuracy : 0.7053
30             95% CI : (0.6863, 0.7237)
31     No Information Rate : 0.2351
32     P-Value [Acc > NIR] : < 2.2e-16
33
34             Kappa : 0.6501
35
36 [1] "Accuracy of test data predictions with top 10 features= 0.7053"

```

---

```

1 > # Check for over-fitting. Criteria: Accuracy change from train to test > 25%
2 > mn_model1_isOF=abs((mn_acc1_tr-mn_acc1_tst)/mn_acc1_tr)
3 > mn_model1_isOF=round(mn_model1_isOF,4)
4 > print(paste("Accuracy drop from training data to test data is",mn_model1_isOF*1
5             00,"%"))
6 > if(mn_model1_isOF>0.25) print("Model is over-fitting") else print("Model is not
7             over-fitting")
8
9 [1] "Accuracy drop from training data to test data is 3.82 %"
10 [1] "Model is not over-fitting"

```

---

## 6 Null and Residual Deviance

Deviance is a measure of goodness of fit of a model. A high value indicates a bad fit. Before extracting the residual deviance of the previously generated model, we need to obtain the null deviance since it is not returned by the multinom model summary. This is done by creating a null model with no features. The deviance of the null model would be the null deviance. The output of the below R code shows the null and residual deviance for the logistic regression classifier with top 10 features.

---

```
1 > # Obtaining the null/residual deviance
2 > mn_model_null=multinom(activity~1,trdf,maxit=1000) # Create the null model
3 > mn_summary_null=summary(mn_model_null) # Summary of null model
4 > mn_nulldeviance=round(mn_summary_null$deviance,2) # Obtain the null deviance
5 > print(paste("Null Deviance =",mn_nulldeviance))
6 > mn_summary1=summary(mn_model1) # Get the summary of the top 10 features model
7 > mn_resdeviance1=round(mn_summary1$deviance,2) # Obtain the residual deviance
8 > print(paste("Residual Deviance with top 10 Features =",mn_resdeviance1))
9
10 [1] "Null Deviance = 21409.92"
11 [1] "Residual Deviance with top 10 Features = 8344.26"
```

---

## 7 Eliminating Correlated Features

In Section 3, we identified 516 correlated features in our data set. Eliminating those features would greatly reduce the number of variables and model complexity without much compromise on the goodness of fit. All feature pairs with a correlation coefficient larger than 0.5 were considered to be correlated, and one of the features in that pair was eliminated. This process left us with 46 features out of the total 561 features in the data set. Then, we again checked the correlation matrix to ensure that the remaining 46 features are uncorrelated.

---

```
1 > # Eliminate correlated predictors
2 > cor_index=which(abs(cordata)>0.5 & abs(cordata)!=1, arr.ind = T) # Get the
   indices where the correlation coefficient >0.5 in the correlation matrix
3 > cor_index=cor_index[!duplicated(cbind(pmax(cor_index[,1], cor_index[,2]), pmin(
   cor_index[,1], cor_index[,2]))),] # Retain one of the two predictors
4 > tbl_cor_index=table(cor_index[,1])
5 > cor_attributes=as.numeric(names(tbl_cor_index)) # Correlated attributes
6 > trdf_uncor=trdf[,-cor_attributes] # Remove correlated predictors from train set
7 > tstdf_uncor=tstdf[,-cor_attributes] # Remove correlated predictors from test set
8 > cordata2=cor(trdf_uncor[,~which(names(trdf_uncor=="activity"))])
9 > print(ifelse(any(abs(cordata2[cordata2!=1])>0.5 ), "Correlated Predictors Exist
   ", "No Correlated Predictors"))
10 [1] "No Correlated Predictors"
```

---



After eliminating correlated features, we examine if multicollinearity exists between the remaining features by computing their Variance Inflation Factor (VIF). This can be performed using the `vif()` function from the `car` package in R. However, this function requires a logit model generated by `glm()`. As stated previously, using `glm()` directly is not suitable for our data since it is not binomial; therefore, we employ a one-vs-all approach for each class. Accordingly, we will have 12 `glm` models, and 12 sets of VIFs, each corresponding to a specific class, as can be seen in the output below.

The `glm` model for classes 6, 7 and 8 did not converge; hence, their VIF values are unrealistically high. Looking at the VIF values of feature X41 in classes 1, 3 and 5, we find that they are considered to be high (larger than 5), which implies that this feature introduces collinearity for those specific classes. However, the VIF value for that same feature is low in other classes, such as class 2, 4, 9 and 11. Therefore, it is not safe to assume that feature X41 can be eliminated due to multicollinearity. The way of aggregating VIFs for multiple classes is still not clear at this point, so we are not able to determine which features exhibit multicollinearity.

---

```

1 > for(i in 1:12){ # Loop over all classes
2 +   class_idx=which(trdf_uncor$activity==i) # Get the indices of containing the
      class i
3 +   vif_df=trdf_uncor # Make a copy of the training set
4 +   vif_df$activity[class_idx]=1 # Set class i to 1
5 +   vif_df$activity[-class_idx]=0 # Set other classes to 0
6 +   bi_model=glm(activity~., vif_df, family="binomial") # Train a binomial glm
      model
7 +   vifs=vif(bi_model) # Compute VIFs for class i
8 +   print(paste("VIFs for class", i))
9 +   print(vifs)
10 + }
11
12 [1] "VIFs for class 1"
13      X1      X4      X38      X39      X40      X41      X56      X70
14 3.098208 2.851807 3.400710 2.498406 2.530201 6.312181 6.383775 2.251108
15      X78      X79      X81      X82      X83      X109      X113      X117
16 1.251561 1.298808 1.158059 1.166577 1.137122 1.548071 1.431990 1.389312
17      X118      X119      X120      X121      X123      X158      X159      X160
18 2.585682 2.302533 2.960153 1.410192 1.368086 2.242708 1.949679 2.578537
19      X161      X162      X163      X189      X193      X197      X198      X199
20 1.130071 1.125181 1.101199 1.288766 1.371259 1.375306 2.193864 1.713175
21      X200      X238      X264      X293      X371      X372      X449      X455
22 2.089800 1.154919 1.384957 1.277993 1.412890 1.381042 1.365529 1.616249
23      X459      X525      X538      X551      X555
24 1.476239 1.133926 1.466514 1.098580 3.088010
25
26
27

```

```

28 [1] "VIFs for class 2"
29      X1      X4      X38      X39      X40      X41      X56      X70
30 1.508995 2.775926 2.311853 2.228741 2.228090 3.526166 3.873505 2.902383
31      X78      X79      X81      X82      X83      X109      X113      X117
32 1.440569 1.274553 1.265775 1.264933 1.162891 1.315260 1.474532 1.300563
33      X118      X119      X120      X121      X123      X158      X159      X160
34 2.246913 1.899664 2.077941 1.340376 1.589314 1.831190 1.568546 1.942967
35      X161      X162      X163      X189      X193      X197      X198      X199
36 1.167426 1.127250 1.068052 1.325442 1.382098 1.389145 1.780115 1.709959
37      X200      X238      X264      X293      X371      X372      X449      X455
38 1.714651 1.239904 1.157375 1.802002 1.289019 1.354888 1.316953 1.507006
39      X459      X525      X538      X551      X555
40 1.893769 1.121596 1.282707 1.125824 1.271238
41
42 [1] "VIFs for class 3"
43      X1      X4      X38      X39      X40      X41      X56      X70
44 3.197803 3.163519 5.892273 4.074426 6.033498 5.243117 2.624277 2.494871
45      X78      X79      X81      X82      X83      X109      X113      X117
46 1.528612 1.384229 1.426188 1.419619 1.187775 1.403864 1.326794 1.511904
47      X118      X119      X120      X121      X123      X158      X159      X160
48 3.843335 3.107188 3.857123 1.309814 1.404070 2.199160 1.846400 1.923367
49      X161      X162      X163      X189      X193      X197      X198      X199
50 1.238444 1.281203 1.139084 1.326498 1.600055 1.622519 2.429850 2.954249
51      X200      X238      X264      X293      X371      X372      X449      X455
52 2.376270 1.168124 1.658927 1.473465 1.588861 1.778404 1.437463 1.800269
53      X459      X525      X538      X551      X555
54 2.289268 1.738784 1.565326 1.282566 1.518434
55
56 [1] "VIFs for class 4"
57      X1      X4      X38      X39      X40      X41      X56      X70
58 1.242305 2.060340 2.338110 2.423989 1.737524 1.636163 1.679673 2.115464
59      X78      X79      X81      X82      X83      X109      X113      X117
60 1.852146 1.970812 1.449579 1.688900 1.437285 1.215422 1.367005 1.331116
61      X118      X119      X120      X121      X123      X158      X159      X160
62 1.657477 1.515006 1.692224 1.866466 2.006416 1.751115 1.813838 2.094575
63      X161      X162      X163      X189      X193      X197      X198      X199
64 1.239992 1.231104 1.347175 1.219874 1.299089 1.333628 1.697989 1.654711
65      X200      X238      X264      X293      X371      X372      X449      X455
66 1.986268 1.073119 1.082173 1.151606 1.258193 1.281281 1.329734 1.662460
67      X459      X525      X538      X551      X555
68 1.589748 1.096493 1.099045 1.095673 1.046986
69
70
71
72

```

```

73 [1] "VIFs for class 5"
74      X1      X4      X38      X39      X40      X41      X56      X70
75 1.348955 2.136724 2.470447 2.614905 1.844320 8.776145 8.791091 2.047321
76      X78      X79      X81      X82      X83      X109      X113      X117
77 2.001416 2.075944 1.377120 1.476964 1.342536 1.201662 1.291318 1.281900
78      X118      X119      X120      X121      X123      X158      X159      X160
79 1.698694 1.587817 1.841608 1.631727 1.615167 1.805812 2.100549 2.308834
80      X161      X162      X163      X189      X193      X197      X198      X199
81 1.135650 1.354483 1.313808 1.199992 1.286884 1.315696 1.770069 1.834057
82      X200      X238      X264      X293      X371      X372      X449      X455
83 2.125129 1.091365 1.087756 1.156486 1.293846 1.281544 1.307419 1.680455
84      X459      X525      X538      X551      X555
85 1.595824 1.091804 1.116073 1.109597 1.093683
86
87 [1] "VIFs for class 6"
88      X1      X4      X38      X39      X40      X41      X56
89 42.855524 63.364793 125.169583 45.882620 50.781249 30.492810 22.450327
90      X70      X78      X79      X81      X82      X83      X109
91 46.825818 33.324388 28.586739 46.463162 31.653346 22.274391 34.813819
92      X113      X117      X118      X119      X120      X121      X123
93 19.535602 28.472943 39.410330 18.021598 87.117388 27.004923 45.886503
94      X158      X159      X160      X161      X162      X163      X189
95 39.991379 13.486358 46.254008 14.506555 44.412735 50.092684 29.738182
96      X193      X197      X198      X199      X200      X238      X264
97 18.123646 28.929756 14.077222 49.631339 100.748304 14.893845 18.278545
98      X293      X371      X372      X449      X455      X459      X525
99 7.112223 59.680747 13.854508 59.449782 45.064500 34.073637 12.019765
100      X538      X551      X555
101 25.336056 14.602903 17.615826
102
103 [1] "VIFs for class 7"
104      X1      X4      X38      X39      X40      X41      X56
105 31.868177 177.166625 11.200648 24.476180 66.891437 75.348373 91.972037
106      X70      X78      X79      X81      X82      X83      X109
107 37.434598 16.520406 34.939562 59.713273 124.091446 157.627019 13.879702
108      X113      X117      X118      X119      X120      X121      X123
109 9.839189 10.126926 18.686022 11.146988 7.841169 51.729168 48.735080
110      X158      X159      X160      X161      X162      X163      X189
111 85.843175 23.136353 28.950654 89.197371 6.862271 14.192065 21.053913
112      X193      X197      X198      X199      X200      X238      X264
113 8.128018 21.513301 6.915824 47.481798 12.653119 20.277080 10.861959
114      X293      X371      X372      X449      X455      X459      X525
115 78.881530 4.580368 53.220159 8.676791 4.589110 35.841888 50.066589
116      X538      X551      X555
117 14.422434 2.887353 6.945561

```

```

118 [1] "VIFs for class 8"
119      X1      X4      X38      X39      X40      X41      X56
120 27.014409 98.810205 56.819967 97.742047 80.593316 77.864826 17.942124
121      X70      X78      X79      X81      X82      X83      X109
122 106.048513 17.882086 36.785059 41.419699 109.140793 46.194450 23.017802
123      X113      X117      X118      X119      X120      X121      X123
124 55.582809 109.438277 80.010039 49.557020 85.066878 69.021493 71.171350
125      X158      X159      X160      X161      X162      X163      X189
126 39.914155 49.006268 103.797623 20.654276 70.514073 27.002994 40.543732
127      X193      X197      X198      X199      X200      X238      X264
128 56.096445 23.143960 58.168688 67.916469 21.668761 35.160758 26.544367
129      X293      X371      X372      X449      X455      X459      X525
130 57.375771 44.072659 35.513751 38.275885 27.207302 37.743787 8.584713
131      X538      X551      X555
132 53.587029 6.724509 18.030759
133
134 [1] "VIFs for class 9"
135      X1      X4      X38      X39      X40      X41      X56      X70
136 6.948660 3.537402 2.900029 3.420179 4.023964 2.290510 2.370297 3.017226
137      X78      X79      X81      X82      X83      X109      X113      X117
138 2.358295 2.266026 2.667164 4.143646 1.829464 1.693695 2.368648 2.031817
139      X118      X119      X120      X121      X123      X158      X159      X160
140 1.857989 2.198627 1.781480 2.294575 4.852022 2.145215 2.154618 2.202328
141      X161      X162      X163      X189      X193      X197      X198      X199
142 1.910851 1.810909 2.381149 2.049078 1.929765 1.772766 1.971122 1.659192
143      X200      X238      X264      X293      X371      X372      X449      X455
144 1.816320 1.851859 1.829741 2.130826 1.674056 1.686078 1.534783 1.988021
145      X459      X525      X538      X551      X555
146 1.923318 1.546294 2.148793 1.446126 2.439367
147
148 [1] "VIFs for class 10"
149      X1      X4      X38      X39      X40      X41      X56
150 41.119522 5.047992 5.441009 8.698068 10.897700 8.434156 8.658131
151      X70      X78      X79      X81      X82      X83      X109
152 5.298329 3.656197 5.441499 8.801977 14.764426 5.966306 7.448855
153      X113      X117      X118      X119      X120      X121      X123
154 5.919016 4.123717 5.709605 4.159300 3.610440 8.415053 14.596603
155      X158      X159      X160      X161      X162      X163      X189
156 3.171314 3.190870 5.292425 3.758870 3.434037 4.750743 4.135214
157      X193      X197      X198      X199      X200      X238      X264
158 5.025502 3.607971 5.664477 3.319357 4.562960 2.690820 2.856199
159      X293      X371      X372      X449      X455      X459      X525
160 2.976780 4.801040 3.994265 3.342047 3.883703 3.859752 3.848630
161      X538      X551      X555
162 3.243979 2.786207 4.630591

```

```

163
164 [1] "VIFs for class 11"
165      X1      X4      X38      X39      X40      X41      X56      X70
166 4.439258 3.868403 3.562122 2.507171 2.224752 2.660717 2.741601 2.777117
167      X78      X79      X81      X82      X83      X109      X113      X117
168 4.261607 2.406433 2.044543 3.795130 1.848667 1.775108 2.117882 1.959908
169      X118      X119      X120      X121      X123      X158      X159      X160
170 1.815609 1.655642 1.609147 2.056802 3.938536 1.694040 1.734767 2.271342
171      X161      X162      X163      X189      X193      X197      X198      X199
172 2.046318 1.502938 1.758021 1.680993 1.814053 1.828079 1.773085 1.551850
173      X200      X238      X264      X293      X371      X372      X449      X455
174 1.841509 1.489378 1.428399 1.606238 1.422491 1.468682 1.308116 1.779544
175      X459      X525      X538      X551      X555
176 1.855883 1.276256 1.436676 1.554791 1.800413
177
178 [1] "VIFs for class 12"
179      X1      X4      X38      X39      X40      X41      X56
180 9.424558 4.998048 6.333884 6.988327 10.844274 4.296336 2.693521
181      X70      X78      X79      X81      X82      X83      X109
182 2.670771 2.740794 3.676265 7.841237 4.741249 7.139081 2.978857
183      X113      X117      X118      X119      X120      X121      X123
184 2.213461 3.388556 2.173051 2.081529 2.926680 3.001478 5.728117
185      X158      X159      X160      X161      X162      X163      X189
186 2.702013 2.177165 3.279492 2.536258 1.856380 3.151271 2.122427
187      X193      X197      X198      X199      X200      X238      X264
188 1.988733 2.259901 2.043928 3.062653 4.949597 2.464531 1.705153
189      X293      X371      X372      X449      X455      X459      X525
190 3.210605 2.653512 2.240503 2.053220 2.091338 2.791467 1.664558
191      X538      X551      X555
192 1.591308 1.476650 2.402801

```

---

## 8 Logistic Regression Model Using Uncorrelated Features

Using the uncorrelated features set, a logistic regression classifier was trained, and the classification accuracy was obtained for the learning and generalization phases, which were found to be 93.5% and 89%, respectively. The drop in accuracy from learning phase to generalization phase was 4.9%; hence, the model is not over-fitting. The residual deviance for this model is 1625.5, which is considerably lower than the one evaluated from the model with top 10 features in section 5, which was 8344.3. Moreover, comparing these two models, there is a significant improvement in the classification accuracy in this model that is trained with uncorrelated features.

---

```

1 > # Create a logistic regression model using uncorrelated features
2 > mn_model2=multinom(formstr,trdf_uncor,maxit=1000) # Train the model
3 > # Obtaining the residual deviance
4 > mn_summary2=summary(mn_model2) # Get the summary of model parameters
5 > mn_resdeviance2=round(mn_summary2$deviance,2) # Obtain the residual deviance
6 > print(paste("Residual Deviance with Uncorrelated Features =",mn_resdeviance2))
7
8 [1] "Residual Deviance with Uncorrelated Features = 1625.51"

```

---

```

1 > # Predict using the train data (Learning Phase)
2 > mn_pred2_tr=predict(mn_model2,trdf_uncor[, -which(names(trdf_uncor)=="activity"
    )], type="class")
3 > mn_cfm2_tr=confusionMatrix(table(trdf_uncor[, which(names(trdf_uncor)=="
    activity")], mn_pred2_tr)) # Confusion Matrix for train data
4 > print("Learning Phase Confusion Matrix")
5 > mn_cfm2_tr
6 > mn_acc2_tr=round(mn_cfm2_tr$overall[['Accuracy']],4) # Accuracy of predictions
    with train data
7 > print(paste("Classification accuracy of learning phase using uncorrelated
    features =",mn_acc2_tr))
8
9 [1] "Learning Phase Confusion Matrix"
10 Confusion Matrix and Statistics
11
12     mn_pred2_tr
13      1    2    3    4    5    6    7    8    9   10   11   12
14 1   824    6    7    0    0    0    0    0    0    0    0    0
15 2     8   730   14    0    0    0    0    0    0    0    0    0
16 3     4    13  663    0    0    0    0    0    0    0    0    0
17 4     0    0    0   725  167    0    0    0    0    0    0    0
18 5     0    0    0   133  878    0    0    0    0    0    0    0
19 6     0    0    0    0    0  1005    0    0    0    0    0    0
20 7     0    0    0    0    0    0   32    0    0    0    0    0
21 8     0    0    0    0    0    0    0   14    0    0    0    0
22 9     0    0    0    0    0    0    0    0   52    0    0    0
23 10    0    0    0    0    0    0    0    0    0   41    0    0
24 11    0    0    0    0    0    0    0    0    0    0   71    0
25 12    0    0    0    0    0    0    0    0    0    0    0   49
26
27 Overall Statistics
28
29             Accuracy : 0.9352
30             95% CI : (0.9284, 0.9416)
31             No Information Rate : 0.1922

```

---

```

32 P-Value [Acc > NIR] : < 2.2e-16
33
34 Kappa : 0.9234
35
36 [1] "Classification accuracy of learning phase using uncorrelated features =
    0.9352"

```

---

```

1 > # Predict using the test data (Generalization Phase)
2 > mn_pred2_tst=predict(mn_model2,tstdf_uncor[, -which(names(tstdf_uncor)=="
    activity")], type="class")
3 > mn_cfm2_tst=confusionMatrix(table(tstdf_uncor[, which(names(tstdf_uncor)=="
    activity")], mn_pred2_tst)) # Confusion Matrix for test data
4 > print("Generalization Phase Confusion Matrix")
5 > mn_cfm2_tst
6 > mn_acc2_tst=round(mn_cfm2_tst$overall[['Accuracy']],4) # Accuracy of
    predictions with test data
7 > print(paste("Classification accuracy of generalization phase using uncorrelated
    features =",mn_acc2_tst))
8
9 [1] "Generalization Phase Confusion Matrix"
10 Confusion Matrix and Statistics
11
12 mn_pred2_tst
13      1  2  3  4  5  6  7  8  9 10 11 12
14 1  377  3  5  0  1  0  0  1  1  0  1  0
15 2   6 299  7  0  0  0  0  0  2  0  7  0
16 3   5  14 287  0  0  0  0  0  0  1  0  0
17 4   0  0  0  0 306 92  0  2  0  1  0  0
18 5   1  0  0  0  63 346  0  0  0  1  0  1
19 6   0  0  0  0  0  0 401  2  0  0  1  4
20 7   0  0  0  0  1  1  0 10  3  0  0  0
21 8   0  0  0  0  3  0  0  0  3  0  2  0
22 9   0  0  0  0  1  0  0  0  1 16  0  5
23 10  0  0  0  0  0  0  0  0  0 14  0  5
24 11  0  1  0  2  0  0  0  0  8  0  8  0
25 12  0  0  0  0  0  0  0  0  0  0  1  7
26
27 Overall Statistics
28
29 Accuracy : 0.8897
30 95% CI : (0.8763, 0.9022)
31 No Information Rate : 0.1888
32 P-Value [Acc > NIR] : < 2.2e-16
33
34 Kappa : 0.8695

```

```

35
36 [1] "Classification accuracy of generalization phase using uncorrelated features
    = 0.8897"

```

---

```

1 > # Check for over-fitting. Criteria: Accuracy change from train to test > 25%
2 > mn_model2_isOF=abs((mn_acc2_tr-mn_acc2_tst)/mn_acc2_tr)
3 > mn_model2_isOF=round(mn_model2_isOF,4)
4 > print(paste("Accuracy drop from training data to test data is",mn_model2_isOF*1
    00,"%"))
5 > if(mn_model2_isOF>0.25) print("Model is over-fitting") else print("Model is not
    over-fitting")
6
7 [1] "Accuracy drop from training data to test data is 4.87 %"
8 [1] "Model is not over-fitting"

```

---

## 9 Logistic Regression Model Performance Metrics

As previously described, the model trained with uncorrelated features has given generally better results than the one trained with the top 10 features from `varImp()`. The performance metric used to compare between these two models was the overall classification accuracy. Here, we will extract other performance metrics, including sensitivity, precision, specificity, recall and balanced accuracy, for each class of the logistic regression model with uncorrelated features. These are obtained from the statistics of the confusion matrices. Subsequently, macro averaging is performed to get the overall metric for all classes, which is computed by simply taking the mean over all classes, as follows:

$$P_{macro} = \frac{1}{N} \sum_{i=1}^N P_i$$

where,  $P$  is the performance metric,  $i$  is an index indicating class number, and  $N$  is the total number of classes, which is 12 in our case. The results for learning and generalization phases are shown in the output of the below R code.

Furthermore, another performance measurement, AUC, was obtained using the `pROC` package in R. This package contains the `multiclass.roc()` function, which computes the ROC curves and AUC for multi-class classification models as defined by *David J. Hand & Robert J. Till*. Since this method implements a One-vs-One approach in generating the ROC curves, we will have 66 curves for our 12 classes. Figures 6 and 7 show the ROC curves for learning and generalization phases, respectively. The plots present ROC curves for 7 different class combinations out of the available 66, and the overall AUCs are provided in the output of the below code.



---

```

1 > # Logistic Regression Performance Parameters
2 > require(pROC)
3 > # Learning Phase
4 > mn_PM2_tr=mn_cfm2_tr$byClass[, c("Balanced Accuracy", "Precision", "Sensitivity",
    "Specificity", "Recall")]
5 > print("Logistic-Regression Learning-Phase Performance Parameters:")
6 > mn_PM2_tr
7 > mn_PMavg2_tr=round(apply(mn_PM2_tr,2,mean),4)
8 > print("Macro Averages:")
9 > t(mn_PMavg2_tr)
10 > mn_prob2_tr=predict(mn_model2, trdf_uncor[, -which(names(trdf_uncor)=="activity"
    )], type="prob")
11 > mn_AUC2_tr=multiclass.roc(trdf_uncor[, which(names(trdf_uncor)=="activity")],
    mn_prob2_tr)
12 > print(paste("Logistic-Regression Learning-Phase AUC:", round(mn_AUC2_tr$auc,4)))
13 > # ROC curves
14 > mn_ROC2_tr=mn_AUC2_tr$rocs
15 > ROC_num=paste("1/", as.character(2), sep="")
16 > plot.roc(mn_ROC2_tr[[ROC_num]][[2]], col=2, main="ROC curves of 7 One-vs-One
    class combinations (Learning Phase)")
17 > for(i in 3:8) {ROC_num=paste("1/", as.character(i), sep="")
18 + lines.roc(mn_ROC2_tr[[ROC_num]][[2]], col=i)}
19 > legend("bottom", legend=c('1/2', '1/3', '1/4', '1/5', '1/6', '1/7', '1/8'), col=2:8,
    lwd=2)
20
21 [1] "Logistic-Regression Learning-Phase Performance Parameters:"
22     Balanced Accuracy Precision Sensitivity Specificity Recall
23 Class: 1           0.9914099 0.9844683 0.9856459 0.9971739 0.9856459
24 Class: 2           0.9849695 0.9707447 0.9746328 0.9953062 0.9746328
25 Class: 3           0.9828604 0.9750000 0.9692982 0.9964226 0.9692982
26 Class: 4           0.9042548 0.8127803 0.8449883 0.9635212 0.8449883
27 Class: 5           0.9049511 0.8684471 0.8401914 0.9697108 0.8401914
28 Class: 6           1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
29 Class: 7           1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
30 Class: 8           1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
31 Class: 9           1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
32 Class: 10          1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
33 Class: 11          1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
34 Class: 12          1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
35 [1] "Macro Averages:"
36     Balanced Accuracy Precision Sensitivity Specificity Recall
37 [1,]           0.9807    0.9676    0.9679    0.9935 0.9679
38
39 [1] "Logistic-Regression Learning-Phase AUC: 0.9987"

```

---

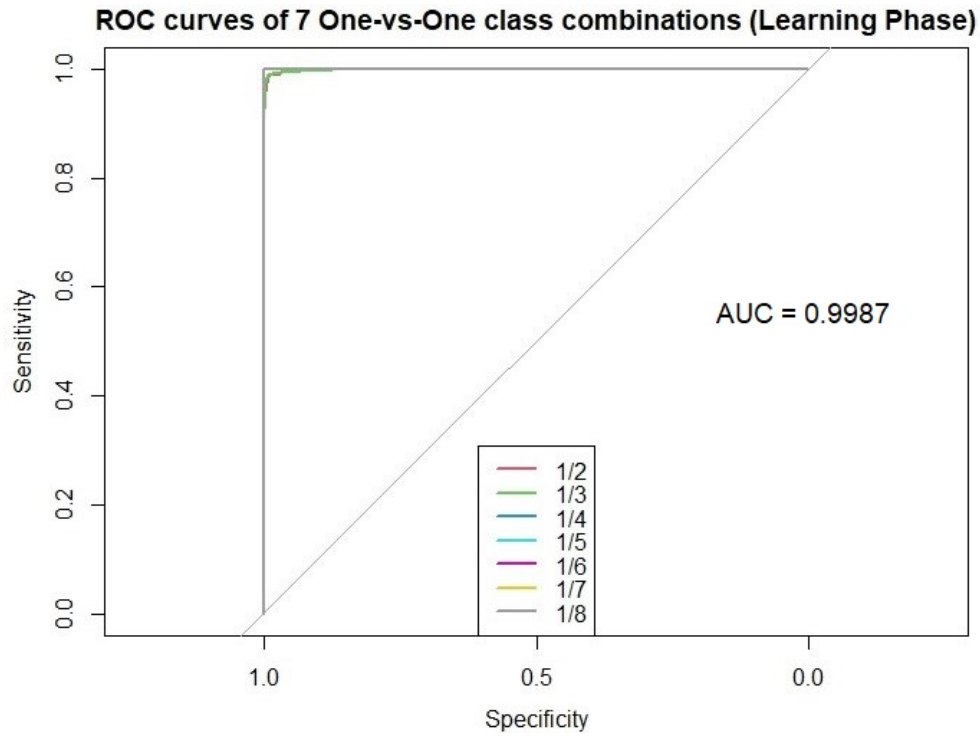
---

```

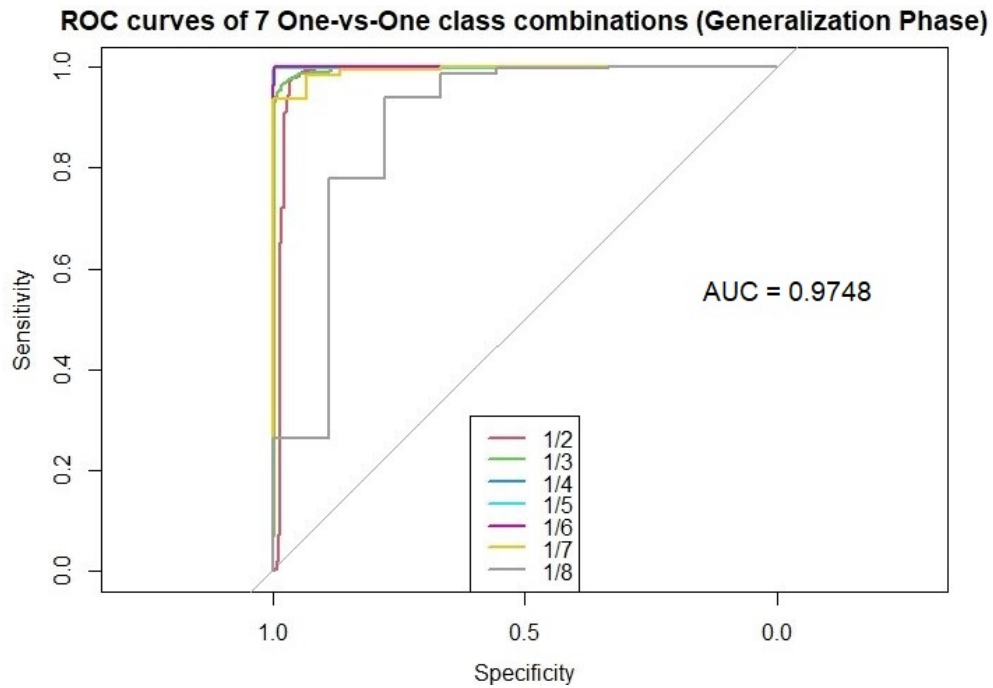
1 > # Generalization Phase
2 > mn_PM2_tst=mn_cfm2_tst$byClass[, c("Balanced Accuracy", "Precision", "
    Sensitivity", "Specificity", "Recall")]
3 > print("Logistic-Regression Generalization-Phase Performance Parameters:")
4 > mn_PM2_tst
5 > mn_PMAvg2_tst=round(apply(mn_PM2_tst,2,mean),4)
6 > print("Macro Averages:")
7 > t(mn_PMAvg2_tst)
8 > mn_prob2_tst=predict(mn_model2, tstdf_uncor[, -which(names(tstdf_uncor)=="
    activity")], type="prob")
9 > mn_AUC2_tst=multiclass.roc(tstdf_uncor[, which(names(tstdf_uncor)=="activity"
    ], mn_prob2_tst)
10 > print(paste("Logistic-Regression Generalization-Phase AUC:", round(mn_AUC2_tst$
    auc,4)))
11 > # ROC curves
12 > mn_ROC2_tst=mn_AUC2_tst$rocs
13 > ROC_num=paste("1/",as.character(2),sep="")
14 > plot.roc(mn_ROC2_tst[[ROC_num]][[2]], col=2, main="ROC curves of 7 One-vs-One
    class combinations (Generalization Phase)")
15 > for(i in 3:8) {ROC_num=paste("1/",as.character(i),sep="")
16 + lines.roc(mn_ROC2_tst[[ROC_num]][[2]],col=i)}
17 > legend("bottom", legend=c('1/2', '1/3', '1/4', '1/5', '1/6', '1/7', '1/8'), col=2:8,
    lwd=2)
18
19 [1] "Logistic-Regression Generalization-Phase Performance Parameters:"
20     Balanced Accuracy Precision Sensitivity Specificity Recall
21 Class: 1           0.9814862 0.9691517 0.9691517 0.9938208 0.9691517
22 Class: 2           0.9661471 0.9314642 0.9432177 0.9890765 0.9432177
23 Class: 3           0.9750119 0.9348534 0.9598662 0.9901575 0.9598662
24 Class: 4           0.8826182 0.7630923 0.8138298 0.9514066 0.8138298
25 Class: 5           0.8757307 0.8398058 0.7863636 0.9650978 0.7863636
26 Class: 6           0.9981865 0.9828431 1.0000000 0.9963731 1.0000000
27 Class: 7           0.8560639 0.6666667 0.7142857 0.9978420 0.7142857
28 Class: 8           0.6862086 0.3333333 0.3750000 0.9974171 0.3750000
29 Class: 9           0.7743417 0.6956522 0.5517241 0.9969592 0.5517241
30 Class: 10          0.8673397 0.7368421 0.7368421 0.9978374 0.7368421
31 Class: 11          0.6514600 0.4210526 0.3076923 0.9952278 0.3076923
32 Class: 12          0.7690151 0.8750000 0.5384615 0.9995686 0.5384615
33
34 [1] "Macro Averages:"
35     Balanced Accuracy Precision Sensitivity Specificity Recall
36 [1,]           0.857    0.7625    0.7247    0.9892 0.7247
37
38 [1] "Logistic-Regression Generalization-Phase AUC: 0.9748"

```

---



**Figure 6:** ROC curves for the learning phase of the logistic regression model.



**Figure 7:** ROC curves for the generalization phase of logistic regression model.

A variance estimation function was created to estimate the variance in the model using 30%, 60% and 100% of the data. The process includes partitioning the data into two subsets, one includes 90% and the other includes 10% of the total number of samples. The 10% subset is reserved for testing, while the 90% subset is placed inside a loop where, in each iteration, a percentage of samples are drawn and used to train a model, and the classification accuracy is computed using the reserved 10% subset. The variance of the accuracies obtained in each iteration would represent the variance estimation for the model. Since we will be using the same function in the subsequent sections, it includes three different model types; Logistic Regression, Naive Bayes, and kNN. The results of variance estimation for the logistic regression model, using 30%, 60% and 100% of the data, are shown in the output below.

---

```

1 > # Variance Estimation
2 > varEst_tridx=sample(1:nrow(trdf_uncor), 0.9*nrow(trdf_uncor), replace=F) # Get
   indices for 90% of the total number of samples
3 > varEst_trdf=trdf_uncor[varEst_tridx,] # Define training data for variance
   estimation
4 > varEst_tstdf=trdf_uncor[-varEst_tridx,] # Define variance estimation partition
5 >
6 > varEst=function(trdf,tstdf,percent,type){
7 +   target_idx=which(names(trdf)=="activity")
8 +   acc_varEstp=c(); # Initialize a variable to store the accuracies computed in
   the loop
9 +   for(i in 1:100){
10 +     varEstp_tridx=sample(1:nrow(trdf), percent/100*nrow(trdf), replace=F) # Take
   samples, percent% of the data
11 +     varEstp_trdf=trdf[varEstp_tridx,]
12 +     if(type=="multinom"){
13 +       mn_model_varEstp=multinom(activity~., varEstp_trdf, maxit=1000, trace=F) #
   Train a multinomial model
14 +       pred_varEstp=predict(mn_model_varEstp, tstdf[,-target_idx], type="class")
   # Predict with variance estimation partition
15 +     }else if(type=="nb"){
16 +       varEstp_trdf$activity=as.factor(varEstp_trdf$activity)
17 +       nb_model_varEstp=naiveBayes(activity~.,data=varEstp_trdf) # Train a naive
   bayes model
18 +       pred_varEstp=predict(nb_model_varEstp, tstdf[,-target_idx], type="class")
   # Predict with variance estimation partition
19 +     }else if(type=="knn"){
20 +       trclass=factor(varEstp_trdf[,target_idx])
21 +       tstclass=factor(tstdf[,target_idx])
22 +       pred_varEstp=knn(varEstp_trdf[,-target_idx], tstdf[,-target_idx], trclass,
   k = 15, prob=TRUE)
23 +     }else{
24 +       print("type should be 'multinom' 'nb' or 'knn'")
25 +       return()

```

```

26 +     }
27 +     u_varEstp=union(pred_varEstp, tstdf[,target_idx]) # Avoids issues when
        number of classes are not equal
28 +     t_varEstp=table(factor(pred_varEstp, u_varEstp), factor(tstdf[,target_idx],
        u_varEstp))
29 +     mn_cfm_varEstp=confusionMatrix(t_varEstp) # Confusion Matrix
30 +     mn_acc_varEstp=mn_cfm_varEstp$overall[['Accuracy']] # Accuracy of
        predictions
31 +     acc_varEstp=c(acc_varEstp,mn_acc_varEstp) # Store
32 + }
33 + mean_varEstp=signif(mean(acc_varEstp),4)
34 + var_varEstp=signif(var(acc_varEstp),4)
35 + varEstp=data.frame(mean_varEstp,var_varEstp)
36 + names(varEstp)=c("Mean of Accuracies","Variance of Accuracies")
37 + return(t(varEstp))
38 + }
39 >
40 > mn_varEst30=varEst(varEst_trdf, varEst_tstdf, 30, type="multinom") # Variance
        estimation using 30% of the data
41 > mn_varEst60=varEst(varEst_trdf, varEst_tstdf, 60, type="multinom") # Variance
        estimation using 60% of the data
42 > mn_varEst100=varEst(varEst_trdf, varEst_tstdf, 100, type="multinom") # Variance
        estimation using 100% of the data
43 >
44 > print("Logistic-Regression Variance Estimation using 30% of data:")
45 > mn_varEst30
46 > print("Logistic-Regression Variance Estimation using 60% of data:")
47 > mn_varEst60
48 > print("Logistic-Regression Variance Estimation using 100% of data:")
49 > mn_varEst100
50
51 [1] "Logistic-Regression Variance Estimation using 30% of data:"
52 Mean of Accuracies    0.83420
53 Variance of Accuracies 0.00016
54
55 [1] "Logistic-Regression Variance Estimation using 60% of data:"
56 Mean of Accuracies    0.8561000
57 Variance of Accuracies 0.0001308
58
59 [1] "Logistic-Regression Variance Estimation using 100% of data:"
60 Mean of Accuracies    0.8842
61 Variance of Accuracies 0.0000

```

---

## 10 Naive Bayes Classifier

Using again the uncorrelated features set, a model is trained based on the Naive Bayes algorithm. The `naiveBayes()` function is used from the `e1071` package, which can generate a multi-class Naive Bayes classifier. The confusion matrices are obtained from predictions in learning and generalization phases, and the performance metrics are computed through the same process that was implemented earlier for the logistic regression model in sections 8 and 9. Confusion matrices and performance metrics are shown in the output below, and the ROC curves are displayed in Fig. 8 and 9. Moreover, the variance estimation function is executed to obtain the variance in the Naive Bayes model.

---

```
1 > # Naive Bayes Performance Parameters
2 > require(e1071)
3 > # Create a Naive Bayes Model
4 > formstr_nb=formula(formstr)
5 > trdf_nb=trdf_uncor
6 > tstdf_nb=tstdf_uncor
7 > trdf_nb$activity=as.factor(trdf_nb$activity) # Makes it run as classification
8 > nb_model1=naiveBayes(formstr_nb,data=trdf_nb) # Train the model
```

---

```
1 > # Predict using train data (Learning Phase)
2 > nb_pred1_tr=predict(nb_model1, trdf_nb[, -which(names(trdf_nb)=="activity")],
  type="class")
3 > nb_cfm1_tr=confusionMatrix(table(trdf_nb[, which(names(trdf_nb)=="activity")],
  nb_pred1_tr)) # Confusion Matrix for train data
4 > print("Naive-Bayes Learning Phase Confusion Matrix")
5 > nb_cfm1_tr
6 > nb_acc1_tr=round(nb_cfm1_tr$overall[['Accuracy']],4) # Accuracy of predictions
  with train data
7 > print(paste("Naive Bayes Learning Phase Accuracy =",nb_acc1_tr))
8 > # Performance Parameters
9 > nb_PM1_tr=nb_cfm1_tr$byClass[, c("Balanced Accuracy", "Precision", "Sensitivity",
  "Specificity", "Recall")]
10 > print("Naive-Bayes Learning-Phase Performance Parameters:")
11 > nb_PM1_tr
12 > nb_PMavg1_tr=round(apply(nb_PM1_tr,2,mean),4)
13 > print("Macro Averages:")
14 > t(nb_PMavg1_tr)
15 > nb_prob1_tr=predict(nb_model1, trdf_nb[, -which(names(trdf_nb)=="activity")],
  type="raw")
16 > nb_AUC1_tr=multiclass.roc(trdf_nb[, which(names(trdf_nb)=="activity")], nb_prob
  1_tr)
17 > print(paste("Naive-Bayes Learning-Phase AUC:", round(nb_AUC1_tr$auc, 4)))
18 > # ROC curves
19 > nb_ROC1_tr=nb_AUC1_tr$rocs
```

---

```

20 > ROC_num=paste("1/",as.character(2),sep="")
21 > plot.roc(nb_ROC1_tr[[ROC_num]][[2]], col=2, main="ROC curves of 7 One-vs-One
      class combinations (Learning Phase)")
22 > for(i in 3:8) {ROC_num=paste("1/",as.character(i),sep="")
23 + lines.roc(nb_ROC1_tr[[ROC_num]][[2]],col=i)}
24 > legend("bottom",legend=c('1/2','1/3','1/4','1/5','1/6','1/7','1/8'), col=2:8,
      lwd=2)
25
26 [1] "Naive-Bayes Learning Phase Confusion Matrix"
27 Confusion Matrix and Statistics
28
29      nb_pred1_tr
30      1  2  3  4  5  6  7  8  9 10 11 12
31 1  766 35 36  0  0  0  0  0  0  0  0
32 2   34 661 55  0  0  0  0  0  0  2  0
33 3   21 34 625  0  0  0  0  0  0  0  0
34 4    0  2  0 393 416  5 16 10 21  1 28  0
35 5    1  5  0 40 934  0 17  0  1  0 13  0
36 6    0  0  0  0 34 910  0  0  2  9 14 36
37 7    0  0  0  0  0  0 30  0  2  0  0  0
38 8    0  0  0  0  0  0  0 14  0  0  0  0
39 9    0  0  0  0  0  0  1  0 40  0 11  0
40 10   0  0  0  0  0  0  0  0  1 27  1 12
41 11   0  0  0  0  0  0  0  0 12  0 59  0
42 12   0  0  0  0  0  1  0  0  1  3  1 43
43
44 Overall Statistics
45
46      Accuracy : 0.8282
47      95% CI : (0.8179, 0.8381)
48      No Information Rate : 0.2546
49      P-Value [Acc > NIR] : < 2.2e-16
50
51      Kappa : 0.7977
52
53 [1] "Naive Bayes Learning Phase Accuracy = 0.8282"
54

```

```

55 [1] "Naive-Bayes Learning-Phase Performance Parameters:"
56       Balanced Accuracy Precision Sensitivity Specificity Recall
57 Class: 1      0.9582428 0.9151732 0.9318735 0.9846121 0.9318735
58 Class: 2      0.9387567 0.8789894 0.8968792 0.9806342 0.8968792
59 Class: 3      0.9306262 0.9191176 0.8729050 0.9883475 0.8729050
60 Class: 4      0.9039405 0.4405830 0.9076212 0.9002598 0.9076212
61 Class: 5      0.8279263 0.9238378 0.6748555 0.9809970 0.6748555
62 Class: 6      0.9862160 0.9054726 0.9934498 0.9789823 0.9934498
63 Class: 7      0.7341888 0.9375000 0.4687500 0.9996277 0.4687500
64 Class: 8      0.7916667 1.0000000 0.5833333 1.0000000 0.5833333
65 Class: 9      0.7488798 0.7692308 0.5000000 0.9977595 0.5000000
66 Class: 10     0.8362027 0.6585366 0.6750000 0.9974055 0.6750000
67 Class: 11     0.7275516 0.8309859 0.4573643 0.9977388 0.4573643
68 Class: 12     0.7357025 0.8775510 0.4725275 0.9988775 0.4725275
69
70 [1] "Macro Averages:"
71       Balanced Accuracy Precision Sensitivity Specificity Recall
72 [1,]      0.8433      0.8381      0.7029      0.9838 0.7029
73
74 [1] "Naive-Bayes Learning-Phase AUC: 0.9898"

```

---

```

1 > # Predict using test data (Generalization Phase)
2 > nb_pred1_tst=predict(nb_model1, tstdf_nb[, -which(names(tstdf_nb)=="activity")
   ], type="class")
3 > nb_cfm1_tst=confusionMatrix(table(tstdf_nb[, which(names(tstdf_nb)=="activity")
   ], nb_pred1_tst)) # Confusion Matrix for test data
4 > print("Naive-Bayes Generalization Phase Confusion Matrix")
5 > nb_cfm1_tst
6 > nb_acc1_tst=round(nb_cfm1_tst$overall[['Accuracy']],4) # Accuracy of
   predictions with test data
7 > print(paste("Naive Bayes Generalization Phase Accuracy =",nb_acc1_tst))
8 > # Check for over-fitting. Criteria: Accuracy change from train to test > 25%
9 > nb_model1_isOF=abs((nb_acc1_tr-nb_acc1_tst)/nb_acc1_tr)
10 > nb_model1_isOF=round(nb_model1_isOF,4)
11 > print(paste("Accuracy drop from training data to test data is",nb_model1_isOF*1
   00,"%"))
12 > if(nb_model1_isOF>0.25) print("Model is over-fitting") else print("Model is not
   over-fitting")
13 > # Performance Parameters
14 > nb_PM1_tst=nb_cfm1_tst$byClass[, c("Balanced Accuracy", "Precision", "
   Sensitivity", "Specificity", "Recall")]
15 > print("Naive-Bayes Generalization-Phase Performance Parameters:")
16 > nb_PM1_tst
17 > nb_PMavg1_tst=round(apply(nb_PM1_tst,2,mean),4)
18 > print("Macro Averages:")

```



```

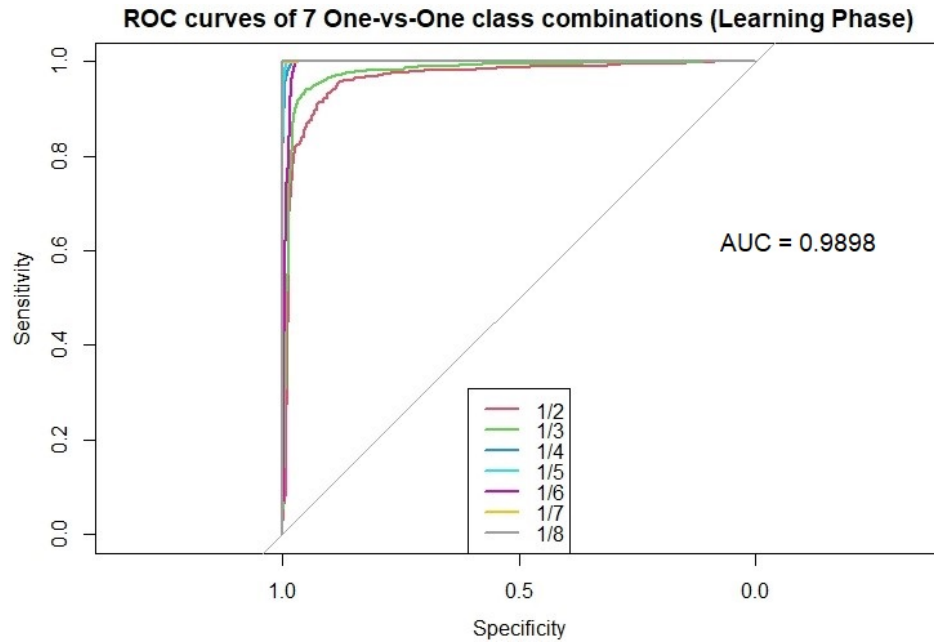
19 > t(nb_PMAvg1_tst)
20 > nb_prob1_tst=predict(nb_model1, tstdf_nb[, -which(names(tstdf_nb)=="activity")
    ], type="raw")
21 > nb_AUC1_tst=multiclass.roc(tstdf_nb[, which(names(tstdf_nb)=="activity")],
    nb_prob1_tst)
22 > print(paste("Naive-Bayes Generalization-Phase AUC:", round(nb_AUC1_tst$auc, 4))
    )
23 > # ROC curves
24 > nb_ROC1_tst=nb_AUC1_tst$rocs
25 > ROC_num=paste("1/",as.character(2),sep="")
26 > plot.roc(nb_ROC1_tst[[ROC_num]][[2]], col=2, main="ROC curves of 7 One-vs-One
    class combinations (Generalization Phase)")
27 > for(i in 3:8) {ROC_num=paste("1/",as.character(i),sep="")
28 + lines.roc(nb_ROC1_tst[[ROC_num]][[2]],col=i)}
29 > legend("bottom",legend=c('1/2','1/3','1/4','1/5','1/6','1/7','1/8'), col=2:8,
    lwd=2)
30
31 [1] "Naive-Bayes Generalization Phase Confusion Matrix"
32 Confusion Matrix and Statistics
33
34     nb_pred1_tst
35      1  2  3  4  5  6  7  8  9 10 11 12
36 1  354 14 21  0  0  0  0  0  0  0  0
37 2   15 282 22  0  0  0  0  0  0  2  0
38 3    8  14 285  0  0  0  0  0  0  0  0
39 4    0  1  0 165 200  6  5  5  8  1 10  0
40 5    0  0  0  20 373  0 11  0  0  0  8  0
41 6    0  0  0  0  12 374  0  0  1  7  5  9
42 7    0  2  0  0  0  0 13  0  0  0  0  0
43 8    0  0  0  0  0  0  0  8  0  0  0  1
44 9    0  0  0  0  0  0  0  0 19  0  4  0
45 10   0  0  0  0  0  0  0  0  0  7  0 12
46 11   0  0  0  0  0  0  0  0  1  0 18  0
47 12   0  0  0  0  0  0  0  0  1  1  0  6
48
49 Overall Statistics
50
51           Accuracy : 0.8168
52           95% CI   : (0.8005, 0.8323)
53     No Information Rate : 0.251
54     P-Value [Acc > NIR] : < 2.2e-16
55
56           Kappa   : 0.7842
57
58 [1] "Naive Bayes Generalization Phase Accuracy = 0.8168"

```

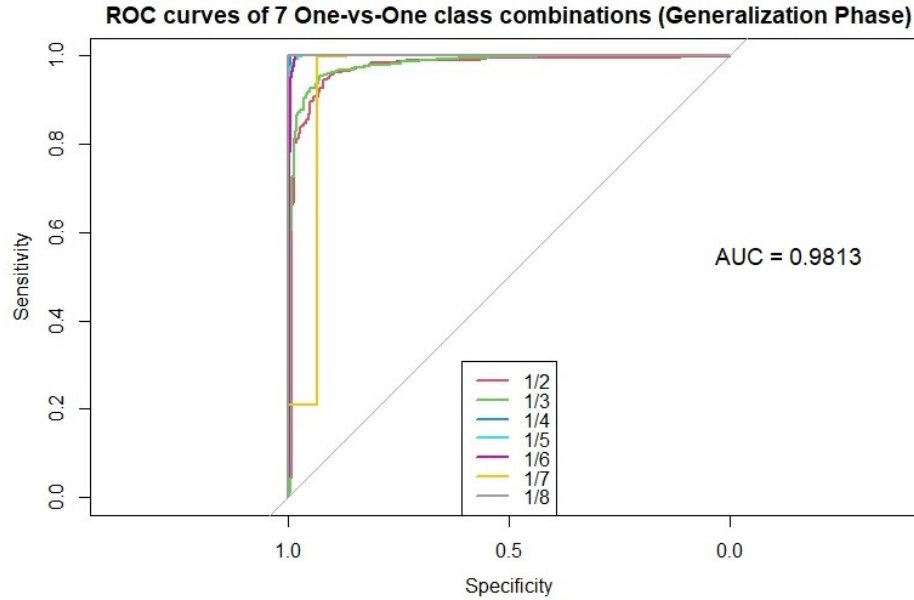
```

59 [1] "Accuracy drop from training data to test data is 1.38 %"
60
61 [1] "Model is not over-fitting"
62
63 [1] "Naive-Bayes Generalization-Phase Performance Parameters:"
64     Balanced Accuracy Precision Sensitivity Specificity Recall
65 Class: 1      0.9605400 0.9100257 0.9389920 0.9820880 0.9389920
66 Class: 2      0.9408162 0.8785047 0.9009585 0.9806739 0.9009585
67 Class: 3      0.9289595 0.9283388 0.8689024 0.9890165 0.8689024
68 Class: 4      0.8909599 0.4114713 0.8918919 0.8900280 0.8918919
69 Class: 5      0.8076350 0.9053398 0.6376068 0.9776632 0.6376068
70 Class: 6      0.9833918 0.9166667 0.9842105 0.9825730 0.9842105
71 Class: 7      0.7237035 0.8666667 0.4482759 0.9991312 0.4482759
72 Class: 8      0.8074766 0.8888889 0.6153846 0.9995686 0.6153846
73 Class: 9      0.8157975 0.8260870 0.6333333 0.9982616 0.6333333
74 Class: 10     0.7161582 0.3684211 0.4375000 0.9948164 0.4375000
75 Class: 11     0.6912704 0.9473684 0.3829787 0.9995622 0.3829787
76 Class: 12     0.6067086 0.7500000 0.2142857 0.9991316 0.2142857
77
78 [1] "Macro Averages:"
79     Balanced Accuracy Precision Sensitivity Specificity Recall
80 [1,]      0.8228      0.7998      0.6629      0.9827 0.6629
81
82 [1] "Naive-Bayes Generalization-Phase AUC: 0.9813"

```



**Figure 8:** ROC curves for the learning phase of the Naive Bayes model.



**Figure 9:** ROC curves for the generalization phase of Naive Bayes model.

---

```

1 > # Variance Estimation
2 > nb_varEst30=varEst(varEst_trdf, varEst_tstdf, 30, type="nb") # Variance
  estimation using 30% of the data
3 > nb_varEst60=varEst(varEst_trdf, varEst_tstdf, 60, type="nb") # Variance
  estimation using 60% of the data
4 > nb_varEst100=varEst(varEst_trdf, varEst_tstdf, 100, type="nb") # Variance
  estimation using 100% of the data
5 > print("Naive-Bayes Variance Estimation using 30% of data:")
6 > nb_varEst30
7 > print("Naive-Bayes Variance Estimation using 60% of data:")
8 > nb_varEst60
9 > print("Naive-Bayes Variance Estimation using 100% of data:")
10 > nb_varEst100
11
12 [1] "Naive-Bayes Variance Estimation using 30% of data:"
13 Mean of Accuracies    0.7876000
14 Variance of Accuracies 0.0001969
15
16 [1] "Naive-Bayes Variance Estimation using 60% of data:"
17 Mean of Accuracies    7.981e-01
18 Variance of Accuracies 7.983e-05
19
20 [1] "Naive-Bayes Variance Estimation using 100% of data:"
21 Mean of Accuracies    0.8051
22 Variance of Accuracies 0.0000

```

---

## 11 kNN Classifier

K-nearest neighbors (kNN) is a non-parametric method for classification, meaning that there are no model parameters to estimate. Therefore, there is no learning phase with this algorithm, since it directly predicts the class of previously unseen data based on the available labeled observations. The `knn()` function from `class` package is used, which takes the training set, its labels, and testing set as inputs, and returns the class predictions for the testing set. Since our number of classes is 12, the value of `k` should be an odd number larger than 12; hence, `k` was set as 15. The confusion matrix, performance metrics, and ROC curve (Fig. 10) are provided below.

```
1 > # kNN Classifier
2 > require(class)
3 > trdf_knn=trdf_uncor[, -which(names(trdf_uncor)=="activity")]
4 > tstdf_knn=tstdf_uncor[, -which(names(trdf_uncor)=="activity")]
5 > trclass_knn=factor(trdf_uncor[, which(names(trdf_uncor)=="activity")])
6 > tstclass_knn=factor(tstdf_uncor[, which(names(tstdf_uncor)=="activity")])
7 > knn_pred1=knn(trdf_knn,tstdf_knn,trclass_knn, k = 15, prob=TRUE)

1 > # Predict using test data (Generalization Phase)
2 > knn_cfm1_tst=confusionMatrix(table(tstclass_knn,knn_pred1)) # Confusion Matrix
   for test data
3 > knn_cfm1_tst
4 > knn_acc1_tst=round(knn_cfm1_tst$overall[['Accuracy']],4) # Accuracy of
   predictions with test data
5 > print(paste("kNN Generalization Phase Accuracy =",knn_acc1_tst))
6 > # Performance Parameters
7 > knn_PM1_tst=knn_cfm1_tst$byClass[, c("Balanced Accuracy", "Precision", "
   Sensitivity", "Specificity", "Recall")]
8 > print("kNN Generalization-Phase Performance Parameters:")
9 > knn_PM1_tst
10 > knn_PMavg1_tst=round(apply(knn_PM1_tst,2,mean),4)
11 > print("Macro Averages:")
12 > t(knn_PMavg1_tst)
13 > knn_prob1_tst=attr(knn_pred1,"prob")
14 > knn_AUC1_tst=multiclass.roc(tstclass_knn, as.ordered(knn_pred1))
15 > print(paste("kNN Generalization-Phase AUC:",round(knn_AUC1_tst$auc,4)))
16 > # ROC curves
17 > knn_ROC1_tst=knn_AUC1_tst$rocs
18 > plot.roc(knn_ROC1_tst[[1]], col=1, main="ROC curves of 7 One-vs-One class
   combinations")
19 > for(i in 2:7) {lines.roc(knn_ROC1_tst[[i]],col=i)}
20 > legend("bottom",legend=c('1/2','1/3','1/4','1/5','1/6','1/7','1/8'), col=1:7,
   lwd=2)
```

21

```

22 Confusion Matrix and Statistics
23
24         knn_pred1
25 tstclass_knn 1  2  3  4  5  6  7  8  9 10 11 12
26         1 383  1  4  0  1  0  0  0  0  0  0
27         2  13 305  3  0  0  0  0  0  0  0  0
28         3  15  5 287  0  0  0  0  0  0  0  0
29         4   0  1  0 249 147  4  0  0  0  0  0
30         5   1  1  0  34 376  0  0  0  0  0  0
31         6   0  0  0  2   2 404  0  0  0  0  0
32         7   0  1  0  0  8   0  6  0  0  0  0
33         8   0  0  0  4   4  0  0  1  0  0  0
34         9   0  1  0  7   1  0  0  0 10  0  4
35        10   0  0  0  2   0  2  0  0  0  4  0
36        11   2  3  0  4   2  0  0  0  3  0  5
37        12   0  1  0  1   0  2  0  0  0  0  4
38
39 Overall Statistics
40
41         Accuracy : 0.8726
42         95% CI : (0.8584, 0.8859)
43         No Information Rate : 0.2321
44         P-Value [Acc > NIR] : < 2.2e-16
45         Kappa : 0.8484
46
47 [1] "kNN Generalization Phase Accuracy = 0.8726"
48
49 [1] "kNN Generalization-Phase Performance Parameters:"
50         Balanced Accuracy Precision Sensitivity Specificity Recall
51 Class: 1          0.9609954 0.9845758 0.9251208 0.9968701 0.9251208
52 Class: 2          0.9740803 0.9501558 0.9561129 0.9920477 0.9561129
53 Class: 3          0.9831861 0.9348534 0.9761905 0.9901816 0.9761905
54 Class: 4          0.8734157 0.6209476 0.8217822 0.9250493 0.8217822
55 Class: 5          0.8374488 0.9126214 0.6950092 0.9798883 0.6950092
56 Class: 6          0.9892491 0.9901961 0.9805825 0.9979156 0.9805825
57 Class: 7          0.9980645 0.4000000 1.0000000 0.9961290 1.0000000
58 Class: 8          0.9982833 0.1111111 1.0000000 0.9965665 1.0000000
59 Class: 9          0.8818112 0.4347826 0.7692308 0.9943917 0.7692308
60 Class: 10         0.9967770 0.2105263 1.0000000 0.9935539 1.0000000
61 Class: 11         0.7747631 0.2631579 0.5555556 0.9939707 0.5555556
62 Class: 12         0.6324698 0.5000000 0.2666667 0.9982729 0.2666667
63 [1] "Macro Averages:"
64         Balanced Accuracy Precision Sensitivity Specificity Recall
65 [1,]          0.9084    0.6094    0.8289    0.9879 0.8289
66 [1] "kNN Generalization-Phase AUC: 0.8216"

```

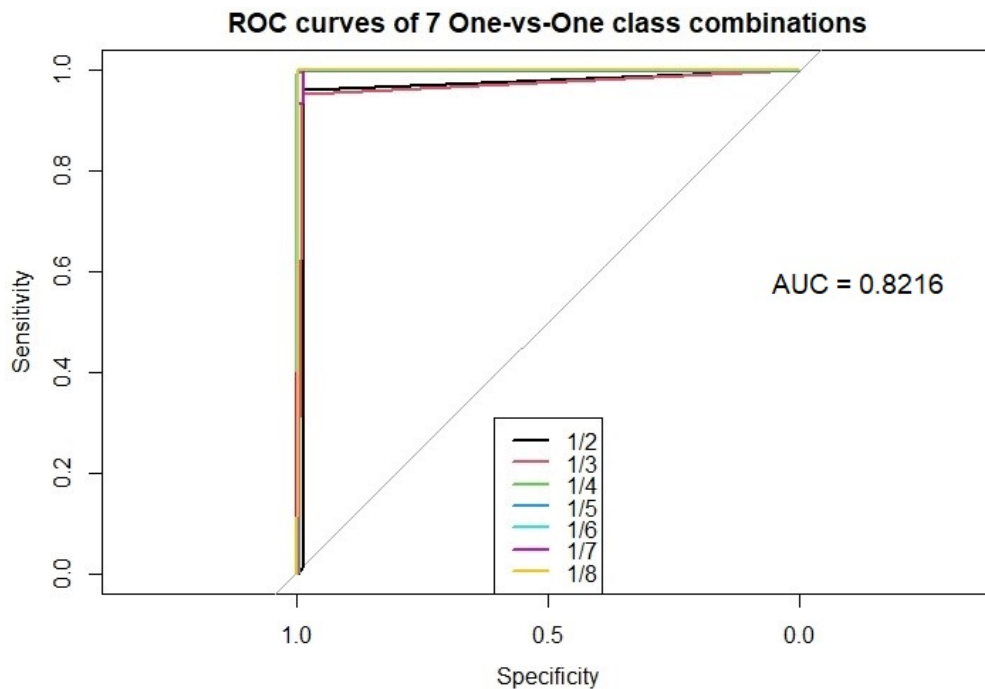
---

```

1 > # Variance Estimation
2 > knn_varEst30=varEst(varEst_trdf, varEst_tstdf, 30, type="knn") # 30% of data
3 > knn_varEst60=varEst(varEst_trdf, varEst_tstdf, 60, type="knn") # 60% of data
4 > knn_varEst100=varEst(varEst_trdf, varEst_tstdf, 100, type="knn") # 100% of data
5 > print("kNN Variance Estimation using 30% of data:")
6 > knn_varEst30
7 > print("kNN Variance Estimation using 60% of data:")
8 > knn_varEst60
9 > print("kNN Variance Estimation using 100% of data:")
10 > knn_varEst100
11
12 [1] "kNN Variance Estimation using 30% of data:"
13 Mean of Accuracies    0.8128000
14 Variance of Accuracies 0.0000908
15
16 [1] "kNN Variance Estimation using 60% of data:"
17 Mean of Accuracies    8.378e-01
18 Variance of Accuracies 5.701e-05
19
20 [1] "kNN Variance Estimation using 100% of data:"
21 Mean of Accuracies    8.585e-01
22 Variance of Accuracies 5.077e-06

```

---



**Figure 10:** ROC curves for the generalization phase of kNN classifier.

## 12 Dealing with the Class Imbalance Issue

Up until this point, we have been working with the data set as it is without addressing the imbalance issue that was observed during EDA in section 3. Out of the three classifiers, it seems that the logistic regression model performed the best in terms of classification accuracy. However, in each of the three classifiers, the classes that suffer imbalance, which are "7", "8", "9", "10", "11", and "12", have low balanced accuracies, as can be seen in the performance metrics table evaluated earlier for each respective classifier. For example, the balanced accuracy of class "11" was 65% in the logistic regression model, 69% in the Naive Bayes model, and 77% in the kNN classifier. Also, for class "12", the accuracies were found to be 77%, 61% and 63% in logistic, Naive Bayes and kNN, respectively. Furthermore, we notice that the precision for these classes, which is measured as  $\frac{TP}{TP+FP}$ , are quite low. These performance metric values are considered to be poor and we need to improve them.

In order to tackle the class imbalance issue, the training data is split into two groups; group A contains the observations with classes from 1 to 6, and group B contains classes 7 to 12, and a classifier is trained for each group. Then, the classification process is divided into two steps. In the first step, we determine the group that the observation belongs to, which is performed using a trained SVM model since it is a non-probabilistic classifier; hence, it is less susceptible to class imbalance issues. The second step is to classify the observation using the model of the group that it belongs to. Moving forward, this process will be loosely referred to as Dual-group Classification.

The below R code implements the dual-group classification process, and the results are provided in the output. Here, we used logistic regression (`multinom()`) to train the two models for groups A and B. Now, we compare the results of this classifier with those obtained for the logistic regression model in section 9. Even though the learning phase accuracy is very slightly worse in the dual-group classifier, the generalization phase accuracy is slightly better. More importantly, there is a noticeable improvement in the balanced accuracies of the imbalanced classes 7 to 12. The accuracy in classes 7, 8, 9, 10, and 11 improved by 10.5%, 1.7%, 12.8%, 8%, and 34.1%, respectively. However, in class 12 it worsened by 2.6%, and in class 8 the accuracy is still undesirably below 70%. Also, the precision values did not demonstrate any improvement.

---

```
1 > ##### Dual-group Classification #####
2 > target = which(names(trdf_uncor)=="activity")
3 > # Change the labels of classes 1-6 to 0, and classes 7-12 to 1 for training
   data
4 > class16_idx=which(trdf_uncor$activity %in% 1:6) # Indeces of rows containing
   classes 1-6
5 > class712_idx=which(trdf_uncor$activity %in% 7:12) # Indeces of rows containing
   classes 7-12
6 > trdf_svm=trdf_uncor # Load the data to a different variable
7 > trdf_svm$activity[class16_idx]=rep(0,length(class16_idx)) # Change the classes
   1-6 to 0s
8 > trdf_svm$activity[class712_idx]=rep(1,length(class712_idx))# Change the classes
```

```

7-12 to 1s
9 > table(trdf_svm$activity)
10
11 0 1
12 5177 259

```

---

```

1 > # Create the model for group detection
2 > trdf_svm$activity=as.factor(trdf_svm$activity) # Makes it run classification
3 > svm_model=svm(activity~.,data=trdf_svm, probability=TRUE) # Train an SVM model
4 > # Partition the training data into two subsets A and B. Subset A contains
    classes 1-6, and subset B contains 7-12
5 > trdfA=trdf_uncor[class16_idx,] # Create subset A
6 > trdfB=trdf_uncor[class712_idx,] # Create subset B
7 > # Train a classifier on subset A
8 > trdfA$activity=as.factor(trdfA$activity)
9 > mn_modelA=multinom(activity~.,trdfA,maxit=1000)
10 > # Train a classifier on subset B
11 > trdfB$activity=as.factor(trdfB$activity)
12 > mn_modelB=multinom(activity~.,trdfB,maxit=1000)

```

---

```

1 > # Predict using full train data
2 > svm_pred_tr = predict(svm_model,trdf_uncor[, -target], type="class") # Predict
    with the group detector
3 > svm_pred_tr_idx0=which(svm_pred_tr==0) # Get indeces of data to be predicted by
    classifier A (classes 1-6)
4 > svm_pred_tr_idx1=which(svm_pred_tr==1) # Get indeces of data to be predicted by
    classifier B (classes 7-12)
5 >
6 > mn_predA_tr = predict(mn_modelA, trdf_uncor[svm_pred_tr_idx0, -target], type="
    class") # Predict the data that belong to subset A with classifier A
7 > mn_predB_tr = predict(mn_modelB, trdf_uncor[svm_pred_tr_idx1, -target], type="
    class") # Predict the data that belong to subset B with classifier B
8 >
9 > grp_pred_tr=rep(0,length(mn_predA_tr)+length(mn_predB_tr)) # Create a vector to
    combine the predictions
10 > grp_pred_tr[svm_pred_tr_idx0]=as.numeric(as.character(mn_predA_tr))
11 > grp_pred_tr[svm_pred_tr_idx1]=as.numeric(as.character(mn_predB_tr))
12 >
13 > grp_cfm_tr=confusionMatrix(table(trdf_uncor[,target], grp_pred_tr)) # Confusion
    Matrix for train data
14 > grp_cfm_tr
15 > print("Dual-group Method Learning-Phase Confusion Matrix")
16 > grp_acc_tr=round(grp_cfm_tr$overall[['Accuracy']],4) # Accuracy of predictions
    with train data

```



```

17 > print(paste("Dual-group Method Learning-Phase Accuracy =", grp_acc_tr))
18 > # Performance Parameters
19 > grp_PM_tr=grp_cfm_tr$byClass[, c("Balanced Accuracy", "Precision", "Sensitivity",
    ", "Specificity", "Recall")]
20 > print("Dual-group Method Learning-Phase Performance Parameters:")
21 > grp_PM_tr
22 > grp_PMavg_tr=round(apply(grp_PM_tr,2,mean),4)
23 > print("Macro Averages:")
24 > t(grp_PMavg_tr)
25
26 [1] "Dual-group Method Learning-Phase Confusion Matrix"
27 Confusion Matrix and Statistics
28
29     grp_pred_tr
30      1    2    3    4    5    6    7    8    9   10   11   12
31  1  824    6    7    0    0    0    0    0    0    0    0    0
32  2    8  730   14    0    0    0    0    0    0    0    0
33  3    4   13  663    0    0    0    0    0    0    0    0
34  4    0    0    0  724  167    0    1    0    0    0    0
35  5    0    0    0  133  878    0    0    0    0    0    0
36  6    0    0    0    0    0 1005    0    0    0    0    0
37  7    0    0    0    1    1    0   30    0    0    0    0
38  8    0    3    0    3    0    0    0    8    0    0    0
39  9    0    0    0    1    0    0    0    0   51    0    0
40 10    0    0    0    0    0    0    0    0    0   41    0
41 11    0    0    0    1    0    0    0    0    0    0   70    0
42 12    0    0    0    0    0    1    0    0    0    0    0   48
43
44 Overall Statistics
45
46             Accuracy : 0.933
47             95% CI : (0.9261, 0.9395)
48     No Information Rate : 0.1924
49     P-Value [Acc > NIR] : < 2.2e-16
50
51             Kappa : 0.9208
52
53 [1] "Dual-group Method Learning-Phase Accuracy = 0.933"
54

```

```

55 [1] "Dual-group Method Learning-Phase Performance Parameters:"
56     Balanced Accuracy Precision Sensitivity Specificity Recall
57 Class: 1      0.9914099 0.9844683 0.9856459 0.9971739 0.9856459
58 Class: 2      0.9830239 0.9707447 0.9707447 0.9953032 0.9707447
59 Class: 3      0.9828604 0.9750000 0.9692982 0.9964226 0.9692982
60 Class: 4      0.9010983 0.8116592 0.8389340 0.9632626 0.8389340
61 Class: 5      0.9045460 0.8684471 0.8393881 0.9697039 0.8393881
62 Class: 6      0.9995030 1.0000000 0.9990060 1.0000000 0.9990060
63 Class: 7      0.9836860 0.9375000 0.9677419 0.9996300 0.9677419
64 Class: 8      0.9994473 0.5714286 1.0000000 0.9988946 1.0000000
65 Class: 9      0.9999071 0.9807692 1.0000000 0.9998143 1.0000000
66 Class: 10     1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
67 Class: 11     0.9999068 0.9859155 1.0000000 0.9998136 1.0000000
68 Class: 12     0.9999072 0.9795918 1.0000000 0.9998144 1.0000000
69
70 [1] "Macro Averages:"
71     Balanced Accuracy Precision Sensitivity Specificity Recall
72 [1,]      0.9788      0.9221      0.9642      0.9933 0.9642

```

---

```

1 > # Predict using test data
2 > svm_pred_tst = predict(svm_model, tstdf_uncor[, -target], type="class") #
    Predict with the group detector
3 > svm_pred_tst_idx0=which(svm_pred_tst==0) # Get indeces of data to be predicted
    by classifier A (classes 1-6)
4 > svm_pred_tst_idx1=which(svm_pred_tst==1) # Get indeces of data to be predicted
    by classifier B (classes 7-12)
5 >
6 > mn_predA_tst = predict(mn_modelA, tstdf_uncor[svm_pred_tst_idx0, -target], type=
    "class") # Predict the data that belong to subset A with classifier A
7 > mn_predB_tst = predict(mn_modelB, tstdf_uncor[svm_pred_tst_idx1, -target], type=
    "class") # Predict the data that belong to subset B with classifier B
8 >
9 > grp_pred_tst=rep(0,length(mn_predA_tst)+length(mn_predB_tst)) # Create a vector
    to combine the predictions
10 > grp_pred_tst[svm_pred_tst_idx0]=as.numeric(as.character(mn_predA_tst))
11 > grp_pred_tst[svm_pred_tst_idx1]=as.numeric(as.character(mn_predB_tst))
12 >
13 > grp_cfm_tst=confusionMatrix(table(tstdf_uncor[, target], grp_pred_tst)) #
    Confusion Matrix for test data
14 > print("Dual-group Method Generalization-Phase Confusion Matrix")
15 > grp_cfm_tst
16 > grp_acc_tst=round(grp_cfm_tst$overall[['Accuracy']],4) # Accuracy of
    predictions with test data
17 > print(paste("Dual-group Method Generalization-Phase Accuracy =", grp_acc_tst))
18 > # Check for over-fitting. Criteria: Accuracy change from train to test > 25%

```

```

19 > grp_model_isOF=abs((grp_acc_tr-grp_acc_tst)/grp_acc_tr)
20 > grp_model_isOF=round(grp_model_isOF,4)
21 > print(paste("Accuracy drop from training data to test data is",grp_model_isOF*1
    00,"%"))
22 > if(grp_model_isOF>0.25) print("Model is over-fitting") else print("Model is not
    over-fitting")
23 > # Performance Parameters
24 > grp_PM_tst=grp_cfm_tst$byClass[, c("Balanced Accuracy", "Precision", "
    Sensitivity", "Specificity", "Recall")]
25 > print("Dual-group Method Learning-Phase Performance Parameters:")
26 > grp_PM_tst
27 > grp_PMavg_tst=round(apply(grp_PM_tst,2,mean),4)
28 > print("Macro Averages:")
29 > t(grp_PMavg_tst)
30
31 [1] "Dual-group Method Generalization-Phase Confusion Matrix"
32 Confusion Matrix and Statistics
33
34     grp_pred_tst
35      1  2  3  4  5  6  7  8  9 10 11 12
36 1  379  4  5  0  1  0  0  0  0  0  0  0
37 2   6 307  8  0  0  0  0  0  0  0  0  0
38 3   5  15 287  0  0  0  0  0  0  0  0  0
39 4   0  0  0 306 93  2  0  0  0  0  0  0
40 5   0  2  0  63 347  0  0  0  0  0  0  0
41 6   0  0  0  0  0 408  0  0  0  0  0  0
42 7   0  2  0  1  2  0  9  1  0  0  0  0
43 8   0  1  0  4  0  1  0  2  0  1  0  0
44 9   0  0  0  1  0  1  1  0 15  1  4  0
45 10  0  0  0  0  0  0  0  0  0 14  0  5
46 11  0  1  0  1  0  0  0  0  5  0 12  0
47 12  0  0  1  0  0  0  0  2  0  0  0  5
48
49 Overall Statistics
50
51             Accuracy : 0.897
52             95% CI : (0.884, 0.9091)
53     No Information Rate : 0.19
54     P-Value [Acc > NIR] : < 2.2e-16
55
56             Kappa : 0.8779
57
58 [1] "Dual-group Method Generalization-Phase Accuracy = 0.897"
59
60 [1] "Accuracy drop from training data to test data is 3.86 %"

```

```

61 [1] "Model is not over-fitting"
62
63 [1] "Dual-group Method Generalization-Phase Performance Parameters:"
64     Balanced Accuracy Precision Sensitivity Specificity Recall
65 Class: 1      0.9833214 0.9742931 0.9717949 0.9948480 0.9717949
66 Class: 2      0.9588476 0.9563863 0.9246988 0.9929965 0.9246988
67 Class: 3      0.9718181 0.9348534 0.9534884 0.9901478 0.9534884
68 Class: 4      0.8826182 0.7630923 0.8138298 0.9514066 0.8138298
69 Class: 5      0.8744339 0.8422330 0.7832957 0.9655720 0.7832957
70 Class: 6      0.9951456 1.0000000 0.9902913 1.0000000 0.9902913
71 Class: 7      0.9487075 0.6000000 0.9000000 0.9974149 0.9000000
72 Class: 8      0.6984953 0.2222222 0.4000000 0.9969905 0.4000000
73 Class: 9      0.8732691 0.6521739 0.7500000 0.9965383 0.7500000
74 Class: 10     0.9364201 0.7368421 0.8750000 0.9978402 0.8750000
75 Class: 11     0.8734881 0.6315789 0.7500000 0.9969762 0.7500000
76 Class: 12     0.7493537 0.6250000 0.5000000 0.9987075 0.5000000
77
78 [1] "Macro Averages:"
79     Balanced Accuracy Precision Sensitivity Specificity Recall
80 [1,]      0.8955      0.7449      0.801      0.99 0.801

```

---

```

1 [1] "Generalization-Phase Performance Parameters of Logistic-Regression Model
   without Dual-group Method, for Comparison:"
2     Balanced Accuracy Precision Sensitivity Specificity Recall
3 Class: 1      0.9814862 0.9691517 0.9691517 0.9938208 0.9691517
4 Class: 2      0.9661471 0.9314642 0.9432177 0.9890765 0.9432177
5 Class: 3      0.9750119 0.9348534 0.9598662 0.9901575 0.9598662
6 Class: 4      0.8826182 0.7630923 0.8138298 0.9514066 0.8138298
7 Class: 5      0.8757307 0.8398058 0.7863636 0.9650978 0.7863636
8 Class: 6      0.9981865 0.9828431 1.0000000 0.9963731 1.0000000
9 Class: 7      0.8560639 0.6666667 0.7142857 0.9978420 0.7142857
10 Class: 8      0.6862086 0.3333333 0.3750000 0.9974171 0.3750000
11 Class: 9      0.7743417 0.6956522 0.5517241 0.9969592 0.5517241
12 Class: 10     0.8673397 0.7368421 0.7368421 0.9978374 0.7368421
13 Class: 11     0.6514600 0.4210526 0.3076923 0.9952278 0.3076923
14 Class: 12     0.7690151 0.8750000 0.5384615 0.9995686 0.5384615
15
16 [1] "Macro Averages:"
17     Balanced Accuracy Precision Sensitivity Specificity Recall
18 [1,]      0.857      0.7625      0.7247      0.9892 0.7247

```

## 13 Dimensionality Reduction Using PCA

Principal Component Analysis (PCA) is a technique used to reduce the dimensionality of the data set and improve model's performance. It involves decomposing the feature columns into principal components (PCs) by solving an eigenvalue problem. The resulting eigenvectors represent the PCs that will be used to train the model, and the eigenvalue of each PC provides information of its importance.

The `prcomp()` function in R is used to carry out the PCA on the entire data set containing all 561 features. A summary of the first 18 PCs is shown in the output, sorted according to their eigenvalues. Since eigenvalues are calculated based on the variance of the data, the standard deviation here corresponds to the square root of the eigenvalue for each PC. Proportion of Variance indicates the amount of variance explained by that single PC, and the Cumulative Proportion is just the amount of variance captured by that PC and the ones before it accumulated. A threshold value of 1.5 was set on the standard deviation, which means that PCs with standard deviation larger than 1.5 will be used to train the model. It was found that there are 28 PCs satisfying this criteria, as shown in Fig. 11, and they account for about 81% of the variance in the data.

---

```
1 > # Perform PCA using all features
2 > target=which(names(trdf)=="activity")
3 > pca=prcomp(trdf[, -target], center = TRUE, scale = TRUE) # Run PCA
4 > pca_summary=summary(pca)
5 > pca_summary$importance[,1:18] # Display a summary of the important principal
    components
6 > pca_std=pca_summary$sdev # Standard deviations of the principal components
7 > high_pca_idx=which(pca_std>=1.5) # Get the principal components with standard
    deviation >= 1.5
8 > print(paste("Number of PCs with standard deviation larger than 1.5 is", length(
    high_pca_idx)))
9 > pca_var_1.5=pca_summary$importance['Cumulative Proportion',min(which(pca_std<1.
    5))] # Amount of variance accounted for in PCs with sdev > 1.5
10 > print(paste("Cumulative variance proportion for PCs with standard deviation
    larger than 1.5 is", pca_var_1.5*100, "%"))
11 > plot(x = 1:40, y = pca_std[1:40], type = "o", xlab = "Principal Component",
    ylab = "Standard Deviation", xaxp=c(1,40,39))
12 > abline(h = 1.5, col="red", lty=5)
13 > legend("top", legend="Standard Deviation = 1.5", col=c("red"), lty=20, bty="n")
14
```

	PC1	PC2	PC3	PC4	PC5	PC6
Standard deviation	16.43764	6.796284	4.23171	3.678691	3.176991	3.088045
Proportion of Variance	0.48163	0.082330	0.03192	0.024120	0.017990	0.017000
Cumulative Proportion	0.48163	0.563970	0.59589	0.620010	0.638000	0.655000

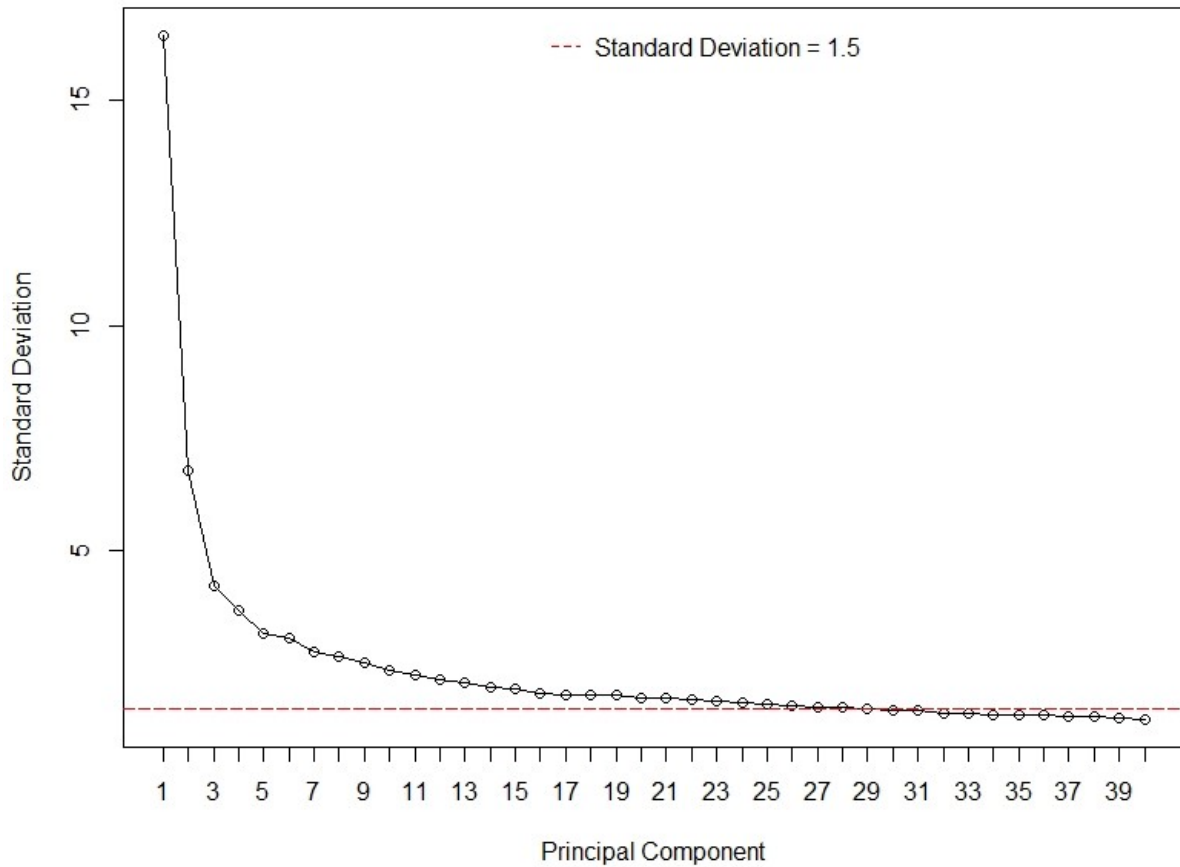
	PC7	PC8	PC9	PC10	PC11	PC12
Standard deviation	2.750629	2.676283	2.52943	2.34298	2.256404	2.15163
Proportion of Variance	0.013490	0.012770	0.01140	0.00979	0.009080	0.00825
Cumulative Proportion	0.668490	0.681250	0.69266	0.70244	0.711520	0.71977

	PC13	PC14	PC15	PC16	PC17	PC18
Standard deviation	2.097646	1.992618	1.94018	1.857391	1.827337	1.819076
Proportion of Variance	0.007840	0.007080	0.00671	0.006150	0.005950	0.005900
Cumulative Proportion	0.727610	0.734690	0.74140	0.747550	0.753500	0.759400

[1] "Number of PCs with standard deviation larger than 1.5 is 28"

[1] "Cumulative variance proportion for PCs with standard deviation larger than 1.5 is 81.225 %"



**Figure 11:** Plot of top 40 PCs and their standard deviations ( $\sqrt{Eigenvalue}$ ).

---

```

1 > # Prepare the training set
2 > trdf_pca=pca$x[,high_pca_idx] # Form a data set from the the principal
    components
3 > trdf_pca=data.frame(trdf_pca) # Convert to data frame type
4 > trdf_pca$activity=trdf$activity # Add the class variable
5 > # Prepare the testing set
6 > tstdf_pca=predict(pca, newdata=tstdf[,-target]) # Get the principal components
    of testing set
7 > tstdf_pca=tstdf_pca[,high_pca_idx] # Form a data set from the the principal
    components
8 > tstdf_pca=data.frame(tstdf_pca) # Convert to data frame type
9 > tstdf_pca$activity=tstdf$activity # Add the class variable
10 > # Train a logistic regression model using principal components
11 > mn_model3=multinom(formstr,trdf_pca,maxit=1000)

```

---

```

1 > # Predict using train data (Learning Phase)
2 > target_pca= which(names(trdf_pca)=="activity")
3 > mn_pred3_tr=predict(mn_model3, trdf_pca[, -target_pca], type="class")
4 > mn_cfm3_tr=confusionMatrix(table(trdf_pca[, target_pca], mn_pred3_tr)) #
    Confusion Matrix for train data
5 > print("PCA Learning Phase Confusion Matrix")
6 > mn_cfm3_tr
7 > mn_acc3_tr=round(mn_cfm3_tr$overall[['Accuracy']],4) # Accuracy of predictions
    with train data
8 > print(paste("PCA Learning Phase Accuracy =",mn_acc3_tr))
9 > # Performance Parameters
10 > mn_PM3_tr=mn_cfm3_tr$byClass[, c("Balanced Accuracy", "Precision", "Sensitivity
    ", "Specificity", "Recall")]
11 > print("PCA Learning-Phase Performance Parameters:")
12 > mn_PM3_tr
13 > mn_PMavg3_tr=round(apply(mn_PM3_tr,2,mean),4)
14 > print("Macro Averages:")
15 > t(mn_PMavg3_tr)
16 > mn_prob3_tr=predict(mn_model3, trdf_pca[, -target_pca], type="probs")
17 > mn_AUC3_tr=multiclass.roc(trdf_pca[, target_pca], mn_prob3_tr)
18 > print(paste("PCA Learning-Phase AUC:", round(mn_AUC3_tr$auc, 4)))
19

```

---

```

20 [1] "PCA Learning Phase Confusion Matrix"
21 Confusion Matrix and Statistics
22
23 mn_pred3_tr
24      1    2    3    4    5    6    7    8    9   10   11   12
25 1  815    8   14    0    0    0    0    0    0    0    0    0
26 2    9  721   22    0    0    0    0    0    0    0    0
27 3   12   23  645    0    0    0    0    0    0    0    0
28 4    0    0    0  733  156    3    0    0    0    0    0
29 5    0    0    0  100  911    0    0    0    0    0    0
30 6    0    0    0    2    0 1003    0    0    0    0    0
31 7    0    0    0    0    0    0   32    0    0    0    0
32 8    0    0    0    0    0    0    0   14    0    0    0
33 9    0    0    0    0    0    0    0    0   39    0   13
34 10   0    0    0    0    0    0    0    0    0   33    0    8
35 11   0    0    0    0    0    0    0    0    9    0   62    0
36 12   0    0    0    0    0    0    0    0    0    6    0   43
37
38 Overall Statistics
39      Accuracy : 0.9292
40      95% CI : (0.922, 0.9359)
41      No Information Rate : 0.1963
42      P-Value [Acc > NIR] : < 2.2e-16
43      Kappa : 0.9162
44
45 [1] "PCA Learning Phase Accuracy = 0.9292"
46
47 [1] "PCA Learning-Phase Performance Parameters:"
48      Balanced Accuracy Precision Sensitivity Specificity Recall
49 Class: 1      0.9850489 0.9737157 0.9748804 0.9952174 0.9748804
50 Class: 2      0.9760792 0.9587766 0.9587766 0.9933817 0.9587766
51 Class: 3      0.9698879 0.9485294 0.9471366 0.9926393 0.9471366
52 Class: 4      0.9216433 0.8217489 0.8778443 0.9654423 0.8778443
53 Class: 5      0.9154536 0.9010880 0.8537957 0.9771115 0.8537957
54 Class: 6      0.9982832 0.9980100 0.9970179 0.9995485 0.9970179
55 Class: 7      1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
56 Class: 8      1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
57 Class: 9      0.9050436 0.7500000 0.8125000 0.9975872 0.8125000
58 Class: 10     0.9223358 0.8048780 0.8461538 0.9985177 0.8461538
59 Class: 11     0.9124939 0.8732394 0.8266667 0.9983212 0.8266667
60 Class: 12     0.9210115 0.8775510 0.8431373 0.9988858 0.8431373
61 [1] "Macro Averages:"
62      Balanced Accuracy Precision Sensitivity Specificity Recall
63 [1,]      0.9523      0.909      0.9115      0.9931 0.9115
64 [1] "PCA Learning-Phase AUC: 0.9964"

```



```

1 > # Predict using test data (Generalization Phase)
2 > mn_pred3_tst=predict(mn_model3, tstdf_pca[, -target_pca], type="class")
3 > mn_cfm3_tst=confusionMatrix(table(tstdf_pca[, target_pca], mn_pred3_tst)) #
    Confusion Matrix for test data
4 > print("PCA Generalization Phase Confusion Matrix")
5 > mn_cfm3_tst
6 > mn_acc3_tst=round(mn_cfm3_tst$overall[['Accuracy']],4) # Accuracy of
    predictions with test data
7 > print(paste("PCA Generalization Phase Accuracy =", mn_acc3_tst))
8 > # Check for over-fitting. Criteria: Accuracy change from train to test > 25%
9 > mn_model3_isOF=abs((mn_acc3_tr-mn_acc3_tst)/mn_acc3_tr)
10 > mn_model3_isOF=round(mn_model3_isOF,4)
11 > print(paste("Accuracy drop from training data to test data is", mn_model3_isOF*
    100,"%"))
12 > if(mn_model3_isOF>0.25) print("Model is over-fitting") else print("Model is not
    over-fitting")
13 > # Performance Parameters
14 > mn_PM3_tst=mn_cfm3_tst$byClass[, c("Balanced Accuracy", "Precision", "
    Sensitivity", "Specificity", "Recall")]
15 > print("PCA Generalization-Phase Performance Parameters:")
16 > mn_PM3_tst
17 > mn_PMavg3_tst=round(apply(mn_PM3_tst,2,mean),4)
18 > print("Macro Averages:")
19 > t(mn_PMavg3_tst)
20 > mn_prob3_tst=predict(mn_model3, tstdf_pca[, -target], type="probs")
21 > mn_AUC3_tst=multiclass.roc(tstdf_pca[, target_pca], mn_prob3_tst)
22 > print(paste("PCA Generalization-Phase AUC:", round(mn_AUC3_tst$auc, 4)))
23
24 [1] "PCA Generalization Phase Confusion Matrix"
25 Confusion Matrix and Statistics
26
27     mn_pred3_tst
28      1  2  3  4  5  6  7  8  9 10 11 12
29 1  379  3  6  0  0  0  0  0  0  0  0  1
30 2   6 303  8  0  0  0  1  2  0  1  0  0
31 3   4  14 287  0  0  0  0  0  1  0  0  1
32 4   0  0  0 295 99  6  1  0  0  0  0  0
33 5   0  0  0  40 370  2  0  0  0  0  0  0
34 6   0  0  0  4  1 400  0  1  0  1  1  0
35 7   1  0  0  0  1  0 11  1  0  0  1  0
36 8   0  0  0  0  0  0  1  6  0  1  0  1
37 9   0  0  0  0  0  0  0  0 11  0 12  0
38 10  0  0  0  0  0  0  0  0  0 13  1  5
39 11  0  0  0  0  0  0  0  0  2  0 17  0
40 12  0  0  0  0  0  0  0  0  0  0  1  7

```

```

41
42 Overall Statistics
43
44         Accuracy : 0.9005
45         95% CI : (0.8876, 0.9123)
46     No Information Rate : 0.2021
47     P-Value [Acc > NIR] : < 2.2e-16
48
49         Kappa : 0.8821
50
51 [1] "PCA Generalization Phase Accuracy = 0.9005"
52
53 [1] "Accuracy drop from training data to test data is 3.09 %"
54
55 [1] "Model is not over-fitting"
56
57 [1] "PCA Generalization-Phase Performance Parameters:"
58     Balanced Accuracy Precision Sensitivity Specificity Recall
59 Class: 1      0.9833214 0.9742931 0.9717949 0.9948480 0.9717949
60 Class: 2      0.9689621 0.9439252 0.9468750 0.9910492 0.9468750
61 Class: 3      0.9718181 0.9348534 0.9534884 0.9901478 0.9534884
62 Class: 4      0.9084968 0.7356608 0.8702065 0.9467871 0.8702065
63 Class: 5      0.8814910 0.8980583 0.7855626 0.9774194 0.7855626
64 Class: 6      0.9881160 0.9803922 0.9803922 0.9958398 0.9803922
65 Class: 7      0.8919940 0.7333333 0.7857143 0.9982736 0.7857143
66 Class: 8      0.7993537 0.6666667 0.6000000 0.9987075 0.6000000
67 Class: 9      0.8902676 0.4782609 0.7857143 0.9948209 0.7857143
68 Class: 10     0.9049541 0.6842105 0.8125000 0.9974082 0.8125000
69 Class: 11     0.7571406 0.8947368 0.5151515 0.9991297 0.5151515
70 Class: 12     0.7331174 0.8750000 0.4666667 0.9995682 0.4666667
71
72 [1] "Macro Averages:"
73     Balanced Accuracy Precision Sensitivity Specificity Recall
74 [1,]      0.8899      0.8166      0.7895      0.9903 0.7895
75
76 [1] "PCA Generalization-Phase AUC: 0.9884"

```

---

---

```

1 [1] "Generalization-Phase Performance Parameters of Logistic-Regression Model
    without PCA, for Comparison:"
2     Balanced Accuracy Precision Sensitivity Specificity Recall
3 Class: 1      0.9814862 0.9691517 0.9691517 0.9938208 0.9691517
4 Class: 2      0.9661471 0.9314642 0.9432177 0.9890765 0.9432177
5 Class: 3      0.9750119 0.9348534 0.9598662 0.9901575 0.9598662
6 Class: 4      0.8826182 0.7630923 0.8138298 0.9514066 0.8138298
7 Class: 5      0.8757307 0.8398058 0.7863636 0.9650978 0.7863636
8 Class: 6      0.9981865 0.9828431 1.0000000 0.9963731 1.0000000
9 Class: 7      0.8560639 0.6666667 0.7142857 0.9978420 0.7142857
10 Class: 8      0.6862086 0.3333333 0.3750000 0.9974171 0.3750000
11 Class: 9      0.7743417 0.6956522 0.5517241 0.9969592 0.5517241
12 Class: 10     0.8673397 0.7368421 0.7368421 0.9978374 0.7368421
13 Class: 11     0.6514600 0.4210526 0.3076923 0.9952278 0.3076923
14 Class: 12     0.7690151 0.8750000 0.5384615 0.9995686 0.5384615
15
16 [1] "Macro Averages:"
17     Balanced Accuracy Precision Sensitivity Specificity Recall
18 [1,]          0.857    0.7625    0.7247    0.9892 0.7247

```

---

Now, we implement the dual-group classification method with PCA and examine if it further improves the imbalance issue. The results below imply that the overall accuracy did not change much, but there is an improvement in the class-specific accuracies of imbalanced classes. Nevertheless, in general, the precision for these classes have worsened.

---

```

1 > ## Dual-group Classification with PCA ###
2 > #####
3 > # Change the labels of classes 1-6 to 0, and classes 7-12 to 1 for training
    data
4 > trdf_svm2=trdf_pca # Load the data to a different variable
5 > trdf_svm2$activity[class16_idx]=rep(0,length(class16_idx)) # Change the classes
    1-6 to 0s
6 > trdf_svm2$activity[class712_idx]=rep(1,length(class712_idx))# Change the
    classes 7-12 to 1s
7 > table(trdf_svm2$activity)
8
9     0     1
10 5177 259

```

---

```

1 > # Create the model for group detection
2 > trdf_svm2$activity=as.factor(trdf_svm2$activity) # Makes it run classification
3 > svm_model2=svm(activity~.,data=trdf_svm2, probability=TRUE) # Train an SVM
    model
4

```

---

```

5 > # Partition the training data into two subsets A and B. Subset A contains
      classes 1-6, and subset B contains 7-12
6 > trdfA2=trdf_pca[class16_idx,] # Create subset A
7 > trdfB2=trdf_pca[class712_idx,] # Create subset B
8 > # Train a classifier on subset A
9 > trdfA2$activity=as.factor(trdfA2$activity)
10 > mn_modelA2=multinom(activity~.,trdfA2,maxit=1000)
11 > # Train a classifier on subset B
12 > trdfB2$activity=as.factor(trdfB2$activity)
13 > mn_modelB2=multinom(activity~.,trdfB2,maxit=1000)



---


1 > # Predict using full train data
2 > svm_pred2_tr = predict(svm_model2,trdf_pca[, -target_pca], type="class") #
      Predict with the group detector
3 > svm_pred2_tr_idx0=which(svm_pred2_tr==0) # Get indices of data to be predicted
      by classifier A (classes 1-6)
4 > svm_pred2_tr_idx1=which(svm_pred2_tr==1) # Get indices of data to be predicted
      by classifier B (classes 7-12)
5 > mn_predA2_tr = predict(mn_modelA2, trdf_pca[svm_pred2_tr_idx0, -target_pca],
      type="class") # Predict the data that belong to subset A with classifier
      A
6 > mn_predB2_tr = predict(mn_modelB2, trdf_pca[svm_pred2_tr_idx1, -target_pca],
      type="class") # Predict the data that belong to subset B with classifier
      B
7 > grp_pred2_tr=rep(0,length(mn_predA2_tr)+length(mn_predB2_tr)) # Create a vector
      to combine the predictions
8 > grp_pred2_tr[svm_pred2_tr_idx0]=as.numeric(as.character(mn_predA2_tr))
9 > grp_pred2_tr[svm_pred2_tr_idx1]=as.numeric(as.character(mn_predB2_tr))
10 >
11 > grp_cfm2_tr=confusionMatrix(table(trdf_pca[,target_pca], grp_pred2_tr)) #
      Confusion Matrix for train data
12 >
13 > print("Dual-group with PCA Learning-Phase Confusion Matrix")
14 > grp_cfm2_tr
15 > grp_acc2_tr=round(grp_cfm2_tr$overall[['Accuracy']],4) # Accuracy of
      predictions with train data
16 > print(paste("Dual-group with PCA Learning-Phase Accuracy =", grp_acc2_tr))
17 > # Performance Parameters
18 > grp_PM2_tr=grp_cfm2_tr$byClass[, c("Balanced Accuracy", "Precision", "
      Sensitivity", "Specificity", "Recall")]
19 > print("Dual-group with PCA Learning-Phase Performance Parameters:")
20 > grp_PM2_tr
21 > grp_PMAvg2_tr=round(apply(grp_PM2_tr,2,mean),4)
22 > print("Macro Averages:")
23 > t(grp_PMAvg2_tr)

```

```

24 [1] "Dual-group with PCA Learning-Phase Confusion Matrix"
25 Confusion Matrix and Statistics
26
27     grp_pred2_tr
28      1    2    3    4    5    6    7    8    9   10   11   12
29  1   815    8   14    0    0    0    0    0    0    0    0    0
30  2     9  721   22    0    0    0    0    0    0    0    0    0
31  3    12   23  645    0    0    0    0    0    0    0    0    0
32  4     0    0    0  733  155    3    1    0    0    0    0    0
33  5     0    0    0   99  912    0    0    0    0    0    0    0
34  6     0    0    0    2    0 1003    0    0    0    0    0    0
35  7     0    3    0    1    1    0   27    0    0    0    0    0
36  8     0    0    0    0    0    0    0   14    0    0    0    0
37  9     0    0    0    0    0    0    0    0   39    0   13    0
38 10     0    0    0    0    0    0    0    0    0   33    0    8
39 11     0    0    0    0    0    1    0    0    9    0   61    0
40 12     0    0    0    0    0    1    0    0    0    6    0   42
41
42 Overall Statistics
43
44             Accuracy : 0.9281
45             95% CI : (0.9209, 0.9348)
46     No Information Rate : 0.1965
47     P-Value [Acc > NIR] : < 2.2e-16
48             Kappa : 0.9149
49
50 [1] "Dual-group with PCA Learning-Phase Accuracy = 0.9281"
51
52 [1] "Dual-group with PCA Learning-Phase Performance Parameters:"
53     Balanced Accuracy Precision Sensitivity Specificity Recall
54 Class: 1      0.9850489 0.9737157 0.9748804 0.9952174 0.9748804
55 Class: 2      0.9741722 0.9587766 0.9549669 0.9933775 0.9549669
56 Class: 3      0.9698879 0.9485294 0.9471366 0.9926393 0.9471366
57 Class: 4      0.9216433 0.8217489 0.8778443 0.9654423 0.8778443
58 Class: 5      0.9156339 0.9020772 0.8539326 0.9773352 0.8539326
59 Class: 6      0.9972940 0.9980100 0.9950397 0.9995483 0.9950397
60 Class: 7      0.9816806 0.8437500 0.9642857 0.9990754 0.9642857
61 Class: 8      1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
62 Class: 9      0.9050436 0.7500000 0.8125000 0.9975872 0.8125000
63 Class: 10     0.9223358 0.8048780 0.8461538 0.9985177 0.8461538
64 Class: 11     0.9112297 0.8591549 0.8243243 0.9981350 0.8243243
65 Class: 12     0.9193502 0.8571429 0.8400000 0.9987003 0.8400000
66 [1] "Macro Averages:"
67     Balanced Accuracy Precision Sensitivity Specificity Recall
68 [1,]      0.9503      0.8931      0.9076      0.993 0.9076

```

---

```

1 > # Predict using test data
2 > svm_pred2_tst = predict(svm_model2, tstdf_pca[, -target_pca], type="class") #
    Predict with the group detector
3 > svm_pred2_tst_idx0=which(svm_pred2_tst==0) # Get indeces of data to be
    predicted by classifier A (classes 1-6)
4 > svm_pred2_tst_idx1=which(svm_pred2_tst==1) # Get indeces of data to be
    predicted by classifier B (classes 7-12)
5 >
6 > mn_predA2_tst = predict(mn_modelA2, tstdf_pca[svm_pred2_tst_idx0, -target_pca],
    type="class") # Predict the data that belong to subset A with classifier
    A
7 > mn_predB2_tst = predict(mn_modelB2, tstdf_pca[svm_pred2_tst_idx1, -target_pca],
    type="class") # Predict the data that belong to subset B with classifier
    B
8 >
9 > grp_pred2_tst=rep(0,length(mn_predA2_tst)+length(mn_predB2_tst)) # Create a
    vector to combine the predictions
10 > grp_pred2_tst[svm_pred2_tst_idx0]=as.numeric(as.character(mn_predA2_tst))
11 > grp_pred2_tst[svm_pred2_tst_idx1]=as.numeric(as.character(mn_predB2_tst))
12 >
13 > grp_cfm2_tst=confusionMatrix(table(tstdf_pca[,target_pca], grp_pred2_tst)) #
    Confusion Matrix for test data
14 > print("Dual-group with PCA Generalization-Phase Confusion Matrix")
15 > grp_cfm2_tst
16 > grp_acc2_tst=round(grp_cfm2_tst$overall[['Accuracy']],4) # Accuracy of
    predictions with test data
17 > print(paste("Dual-group with PCA Generalization-Phase Accuracy =", grp_acc2_tst
    ))
18 > # Check for over-fitting. Criteria: Accuracy change from train to test > 25%
19 > grp_model2_isOF=abs((grp_acc2_tr-grp_acc2_tst)/grp_acc2_tr)
20 > grp_model2_isOF=round(grp_model2_isOF,4)
21 > print(paste("Accuracy drop from training data to test data is",grp_model2_isOF*
    100,"%"))
22 > if(grp_model2_isOF>0.25) print("Model is over-fitting") else print("Model is
    not over-fitting")
23 > # Performance Parameters
24 > grp_PM2_tst=grp_cfm2_tst$byClass[, c("Balanced Accuracy", "Precision", "
    Sensitivity", "Specificity", "Recall")]
25 > print("Dual-group with PCA Generalization-Phase Performance Parameters:")
26 > grp_PM2_tst
27 > grp_PMavg2_tst=round(apply(grp_PM2_tst,2,mean),4)
28 > print("Macro Averages:")
29 > t(grp_PMavg2_tst)
30
31

```

```

32 [1] "Dual-group with PCA Generalization-Phase Confusion Matrix"
33 Confusion Matrix and Statistics
34
35     grp_pred2_tst
36      1  2  3  4  5  6  7  8  9 10 11 12
37 1  379  3  6  0  1  0  0  0  0  0  0  0
38 2   6 307  8  0  0  0  0  0  0  0  0
39 3   4  14 288  0  1  0  0  0  0  0  0
40 4   0  0  0 295 100  6  0  0  0  0  0
41 5   0  0  0  40 370  2  0  0  0  0  0
42 6   0  0  0  3  1 403  0  0  0  0  1  0
43 7   1  0  0  0  4  0  7  3  0  0  0  0
44 8   0  1  0  0  0  0  1  7  0  0  0  0
45 9   0  0  0  0  0  0  2  0 10  0 11  0
46 10  0  0  0  0  0  0  0  0  0 13  1  5
47 11  0  0  0  0  0  0  1  0  2  0 16  0
48 12  0  0  0  0  1  0  0  0  0  0  1  6
49
50 Overall Statistics
51           Accuracy : 0.9013
52           95% CI : (0.8885, 0.9131)
53   No Information Rate : 0.2051
54   P-Value [Acc > NIR] : < 2.2e-16
55           Kappa : 0.883
56
57 [1] "Dual-group with PCA Generalization-Phase Accuracy = 0.9013"
58
59 [1] "Accuracy drop from training data to test data is 2.89 %"
60
61 [1] "Model is not over-fitting"
62
63 [1] "Dual-group with PCA Generalization-Phase Performance Parameters:"
64           Balanced Accuracy Precision Sensitivity Specificity Recall
65 Class: 1      0.9833214 0.9742931 0.9717949 0.9948480 0.9717949
66 Class: 2      0.9688182 0.9563863 0.9446154 0.9930209 0.9446154
67 Class: 3      0.9721391 0.9381107 0.9536424 0.9906358 0.9536424
68 Class: 4      0.9097975 0.7356608 0.8727811 0.9468138 0.8727811
69 Class: 5      0.8756963 0.8980583 0.7740586 0.9773341 0.7740586
70 Class: 6      0.9889656 0.9877451 0.9805353 0.9973958 0.9805353
71 Class: 7      0.8164577 0.4666667 0.6363636 0.9965517 0.6363636
72 Class: 8      0.8495692 0.7777778 0.7000000 0.9991383 0.7000000
73 Class: 9      0.9138637 0.4347826 0.8333333 0.9943941 0.8333333
74 Class: 10     0.9987058 0.6842105 1.0000000 0.9974116 1.0000000
75 Class: 11     0.7660148 0.8421053 0.5333333 0.9986962 0.5333333
76 Class: 12     0.7722962 0.7500000 0.5454545 0.9991379 0.5454545

```

```

77
78 [1] "Macro Averages:"
79     Balanced Accuracy Precision Sensitivity Specificity Recall
80 [1,]          0.9013    0.7871    0.8122    0.9904 0.8122

```

---

```

1 [1] "PCA without Dual-Group Generalization-Phase Performance Parameters for
   Comparison:"
2     Balanced Accuracy Precision Sensitivity Specificity Recall
3 Class: 1          0.9833214 0.9742931 0.9717949 0.9948480 0.9717949
4 Class: 2          0.9689621 0.9439252 0.9468750 0.9910492 0.9468750
5 Class: 3          0.9718181 0.9348534 0.9534884 0.9901478 0.9534884
6 Class: 4          0.9084968 0.7356608 0.8702065 0.9467871 0.8702065
7 Class: 5          0.8814910 0.8980583 0.7855626 0.9774194 0.7855626
8 Class: 6          0.9881160 0.9803922 0.9803922 0.9958398 0.9803922
9 Class: 7          0.8919940 0.7333333 0.7857143 0.9982736 0.7857143
10 Class: 8          0.7993537 0.6666667 0.6000000 0.9987075 0.6000000
11 Class: 9          0.8902676 0.4782609 0.7857143 0.9948209 0.7857143
12 Class: 10         0.9049541 0.6842105 0.8125000 0.9974082 0.8125000
13 Class: 11         0.7571406 0.8947368 0.5151515 0.9991297 0.5151515
14 Class: 12         0.7331174 0.8750000 0.4666667 0.9995682 0.4666667
15
16 [1] "Macro Averages:"
17     Balanced Accuracy Precision Sensitivity Specificity Recall
18 [1,]          0.8899    0.8166    0.7895    0.9903 0.7895
19
20 [1] "PCA Generalization-Phase AUC: 0.9884"

```

---